



**Barcelona  
Supercomputing  
Center**  
*Centro Nacional de Supercomputación*

# COMPSS Tutorial

February 9-10 2016, Valencia

Daniele Lezzi, Carlos Diaz



EXCELENCIA  
SEVERO  
OCHOA

# Outline (Feb 9<sup>th</sup> 2016)

## « Session 1 (15:00 – 16:30): Introduction to COMPSs

- Programming model

- Java Syntax
- Demo: First Java example
- Python syntax
- Demo: First Python example
- Other sample codes

## « Coffee break (16:30 – 16:45)

## « Session 2 (16:45 – 18:15):

- COMPSs execution environments

# Outline (Feb 9<sup>th</sup> 2016)

## « Session 3 (18:15 - 19:45) Hands-on I

- Virtual Machine Setup
- Java Hands-on
  - Word-count taskified code
  - Configuration, monitoring, debugging
  - Graph generation

## « Session 4 (19:45 – 21:00): Hands-on II

- Python Hands-on
  - Word-count without annotations
  - Annotate tasks
  - Execution in MN 3
  - Overview of tracing, trace analysis
  - Code optimization (reduce in a tree)

## « Final notes

# Outline (Feb 10<sup>th</sup> 2016)

## ● Session 5 (10:00 - 12:00) Hands-on III – EGI

- Virtual Machine Setup - Daniele
- Java Hands-on – Daniele



**Barcelona  
Supercomputing  
Center**

*Centro Nacional de Supercomputación*

# Introduction

# Motivation

## « New complex architectures constantly emerging

- With their own way of programming them
  - Fine grain: e.g. APIs to run with GPUs, NVMs (Non-Volatile Memories)
  - Coarse grain: e.g. APIs to deploy in Clouds
- **Difficult** for programmers
  - Higher learning curve / Time To Market (TTM)
  - What about non computer scientists???
- **Difficult** to understand what is going on during execution
  - Was it fast? Could it be even faster? Am I paying more than I should? (**Efficiency**)
- Tune your application for each architecture (or cluster)
  - E.g. partitioning data among nodes

# Motivation

## « Create tools that make user's life **easier**

- Intermediate layer: let the difficult parts to those tools
  - Act on behalf of the user
  - Distributing the work through resources
  - Dealing with architecture specifics
  - Automatically improving performance
- Tools for visualization
  - Monitoring
  - Performance analysis

# The parallel programming revolution

## « Parallel programming in the past

- Where to place data
- What to run, where
- How to communicate

Schedule @ programmers mind

Static

Complexity: Divergence between  
our mental model and reality

Variability

## « Parallel programming in the future

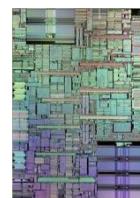
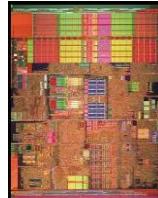
- What do I need to compute
- What data do I need to use
- Hints (not necessarily very precise) on  
potential concurrency, locality,...
- **YOU PROGRAM SEQUENTIALLY!!!**

Schedule @ system

Dynamic

# Living in the programming revolution

« At the beginning there was one language



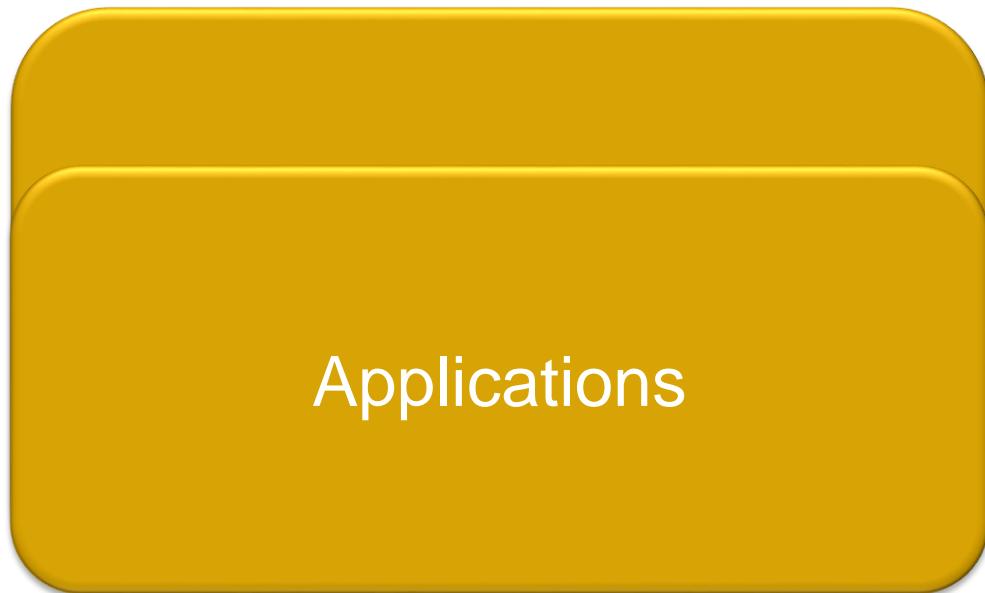
Programs  
“decoupled”  
from hardware

Simple interface  
Sequential program

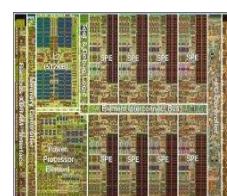
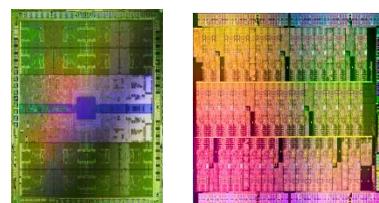
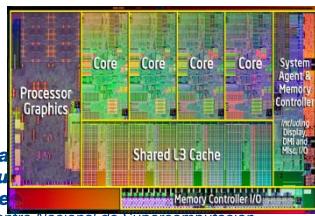
ILP

# Living in the programming revolution

## « Multicores made the interface to leak



Application logic  
+  
**Platform specificities**



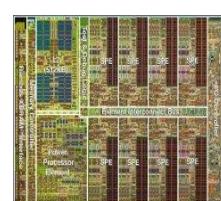
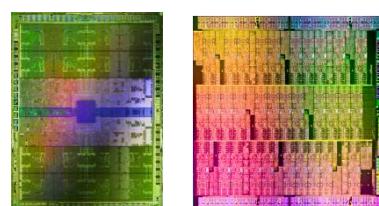
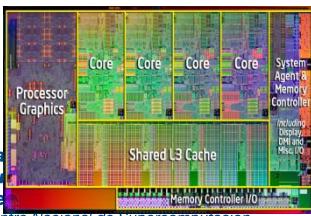
Address spaces  
(hierarchy,  
transfer), control  
flows, ...

# BSC Vision in the programming revolution (StarSs)

## « Need to decouple again



PM: High-level, clean, abstract interface



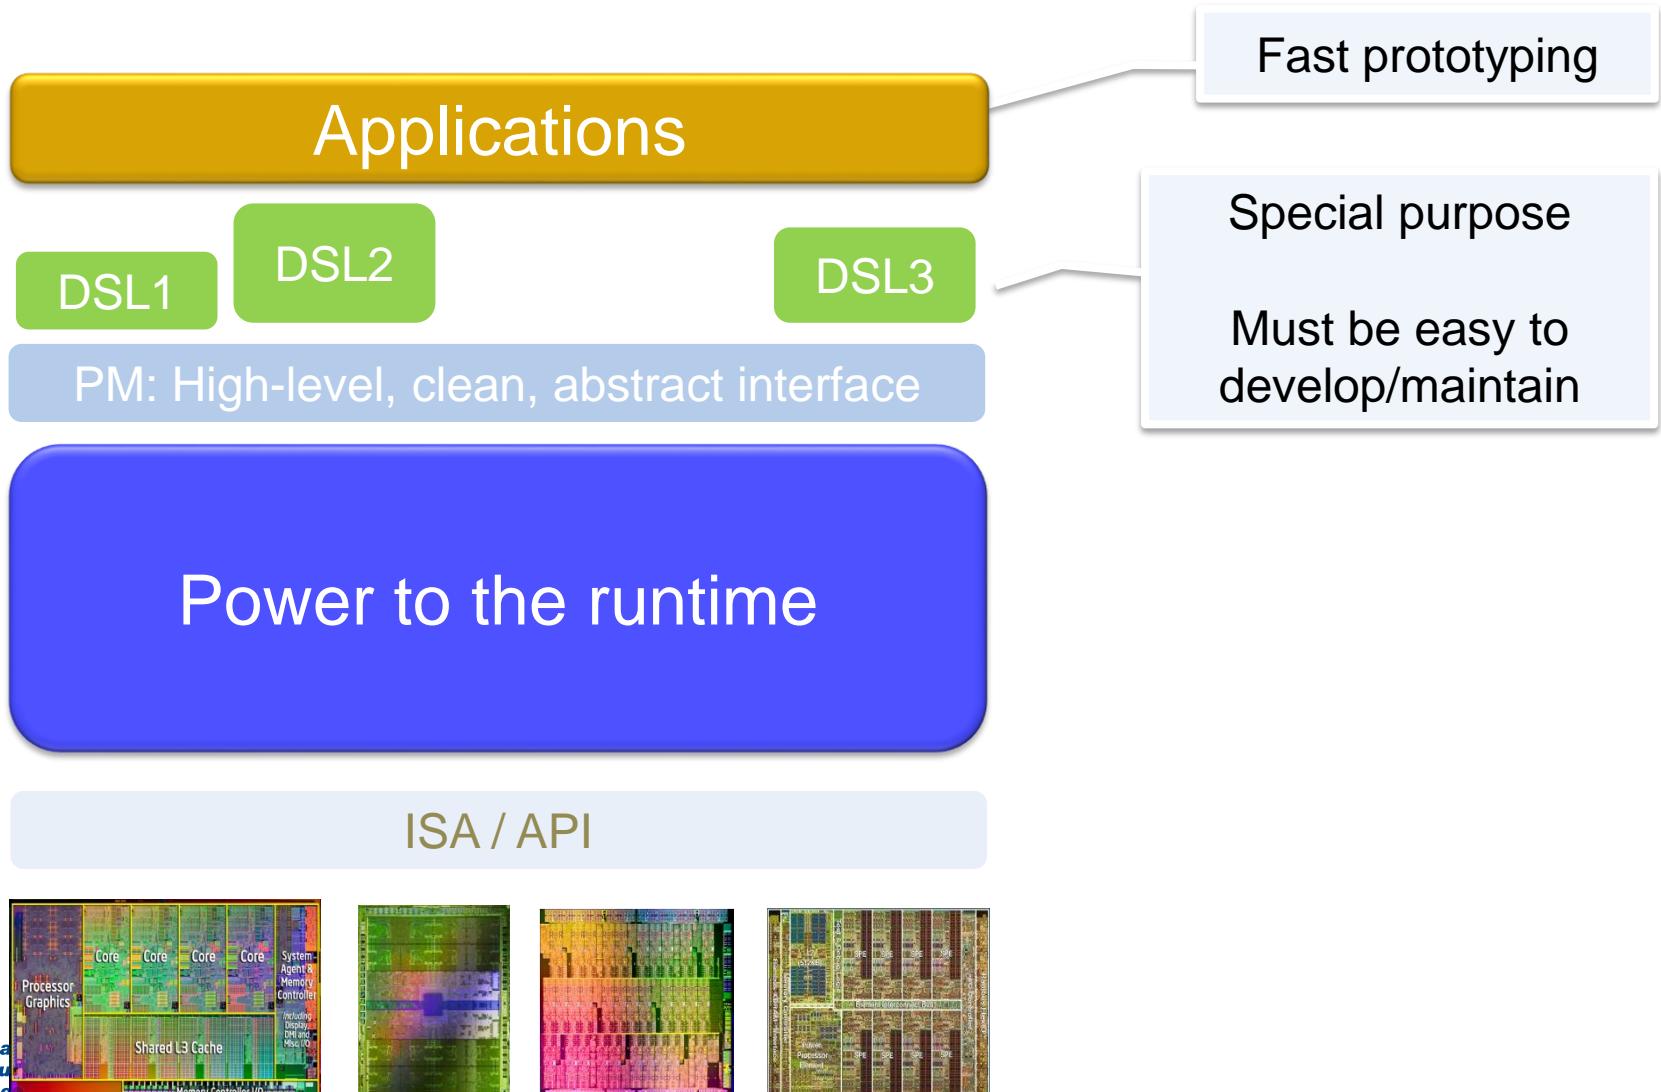
Application logic

Arch. independent

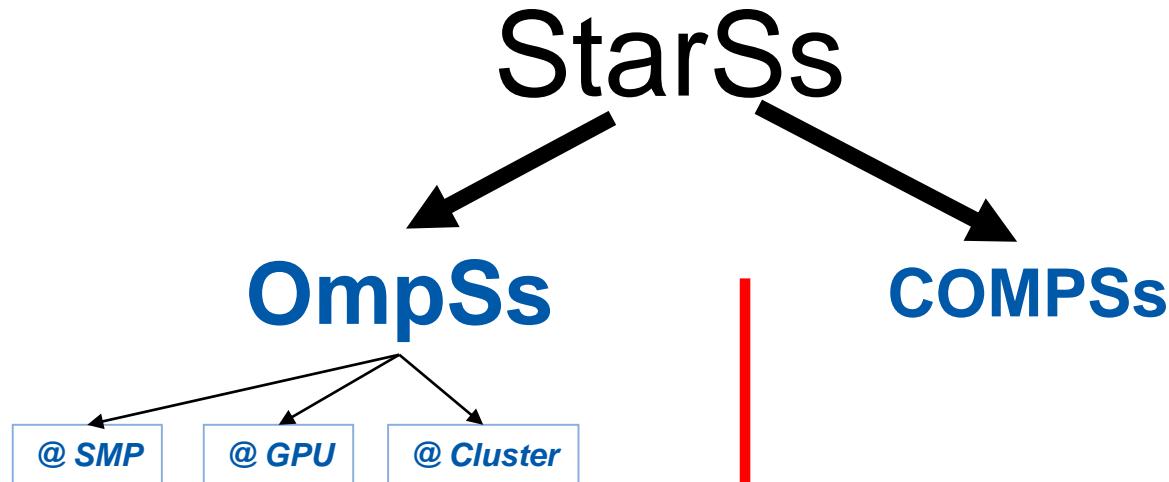
General purpose  
Task based  
Single address space

“Reuse”  
architectural ideas  
under  
new constraints  
(Out-Of-Order  
execution)

# BSC Vision in the programming revolution (StarSs)



# The StarSs “Granularities”



## Average task Granularity:

100 microseconds – 10 milliseconds

10 ms - 1 day

## Address space to compute dependences:

Memory

Files, Objects, NVMs

## Language bindings:

C, C++, FORTRAN

Java, C/C++, Python

SMPs, Clusters

Clusters, Clouds



***Barcelona  
Supercomputing  
Center***

*Centro Nacional de Supercomputación*

**COMPSs**

# Let's narrow the StarSs idea...for Distributed Architectures

## « Cluster / Cloud applications are complex to develop

- Even more if you want to run things in parallel
- **Goal 1: Keep a Sequential Programming Paradigm**
  - Writing an application for a computational distributed infrastructure should be as easy as writing a sequential application
- **Goal 2: Exploit parallelism**
  - Run it as fast as possible

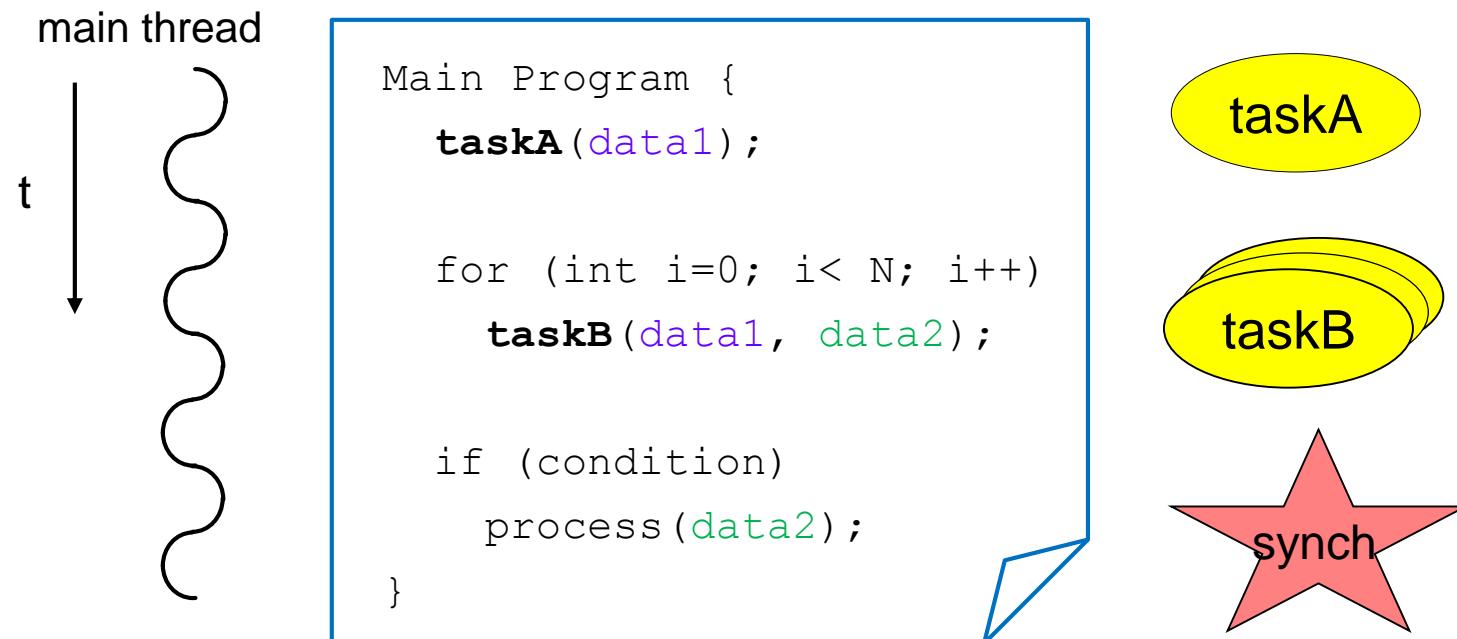
## « Target applications: composed of tasks, most of them repetitive

- Granularity of the tasks: enough to be distributed (simulators, ...)
- Data: files, objects, arrays and primitive types

# Programming Model: Properties (I)

## « Based on sequential programming

- No APIs, no threading, no messaging
- No parallel constructs, no pragmas
- Sequential consistency



# Runtime System

Application

Task Selection Interface

Runtime System



Grid



Cluster



Cloud

# Supported Features

## « Basic Features:

- Data dependency analysis
- Data transfer
- Task scheduling
- Resource management
- Results collection
- Fault tolerance
- Method and Web Service Tasks

## « Advanced Features:

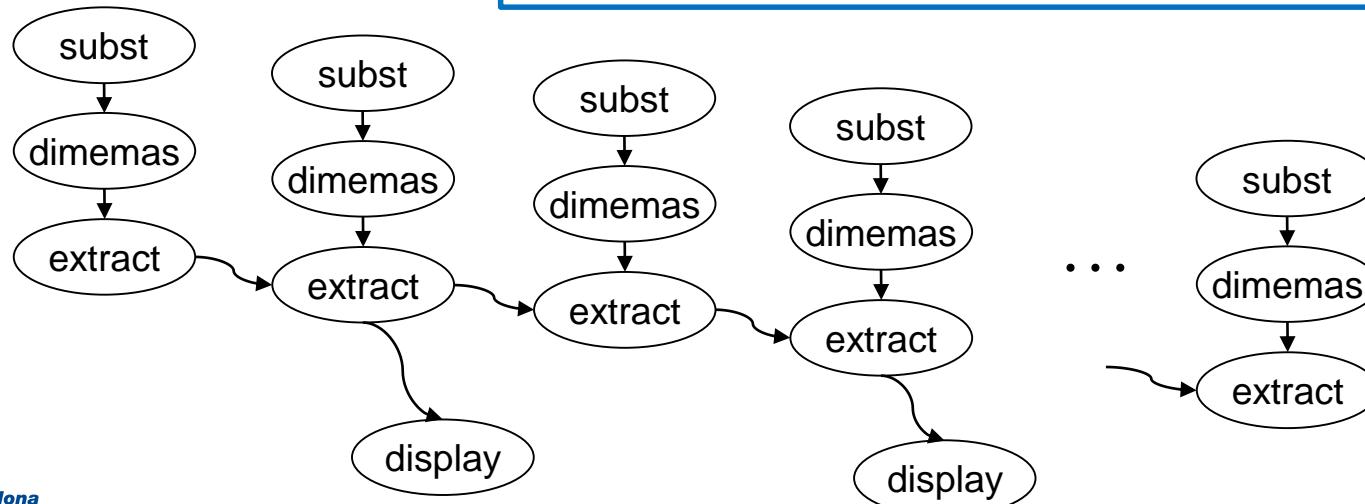
- Shared disks support
- Constraints based scheduling
- Task versioning support

# Programming Model: Dependency detection

## Automatic on-the-fly creation of a task dependency graph

Main Program

```
for (int i = 0; i < N; i++) {  
    newBWD = random();  
    subst(refCFG, newBWD, newCFG);  
    dimemas(newCFG, trace, dimOUT);  
    extract(newBWD, dimOUT, finalOUT);  
    if (i % 2 == 0) display(finalOUT);  
}
```

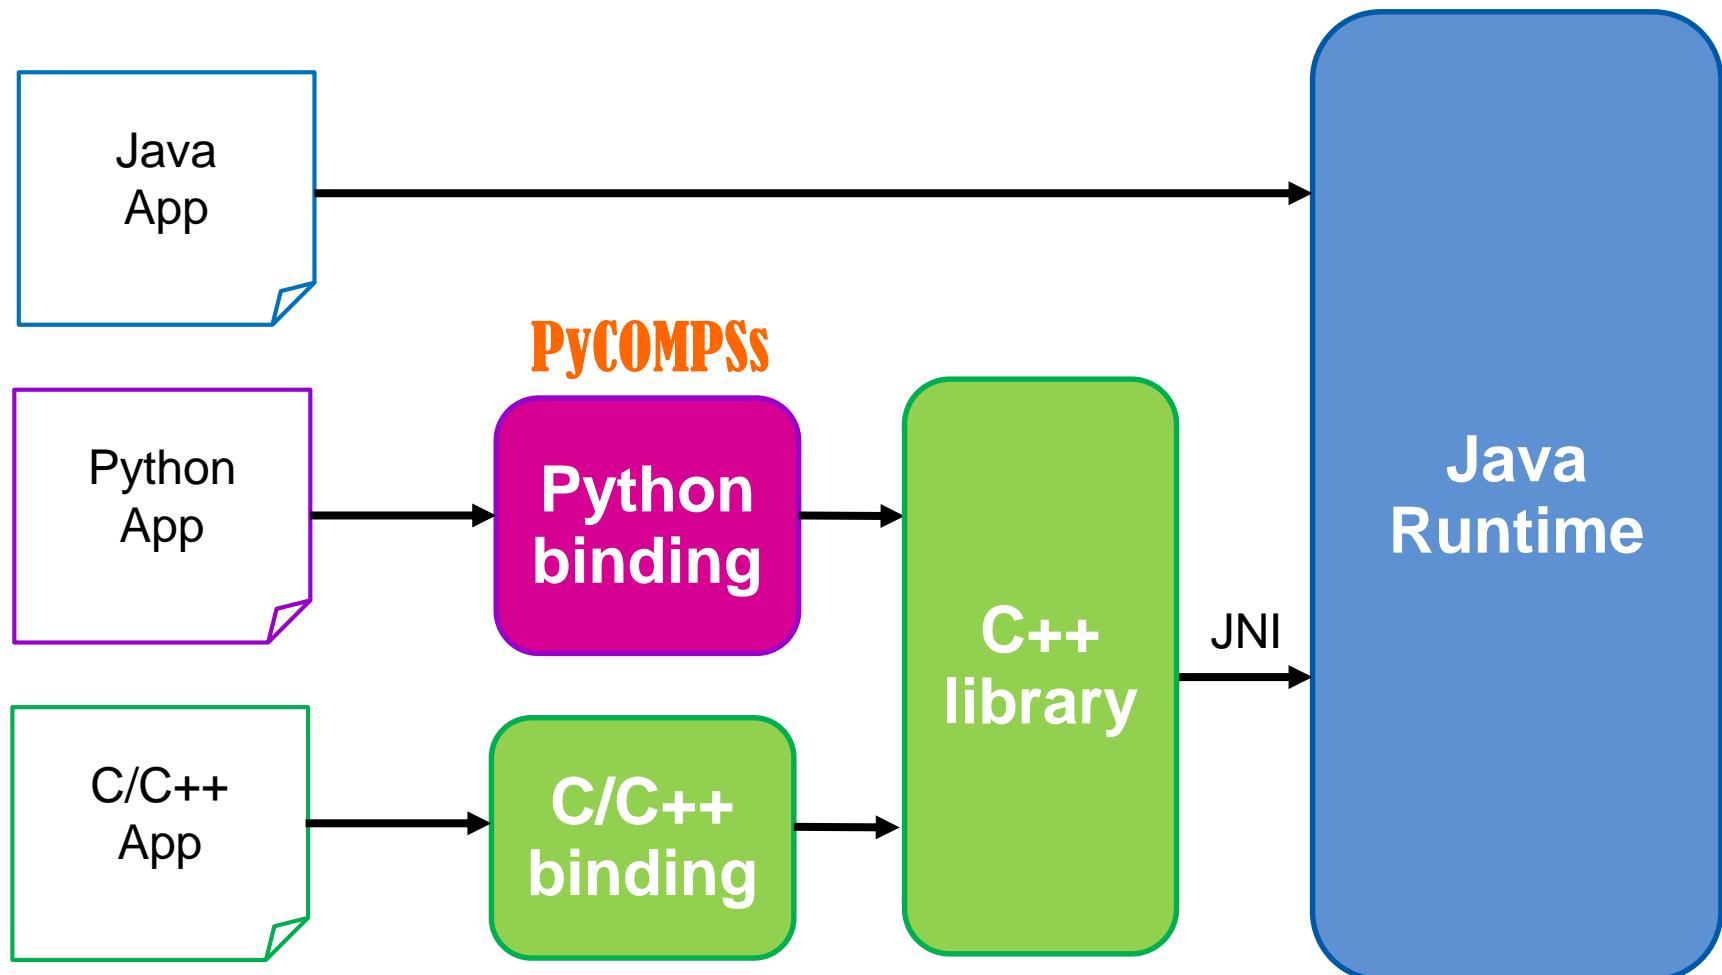




***Barcelona  
Supercomputing  
Center***  
*Centro Nacional de Supercomputación*

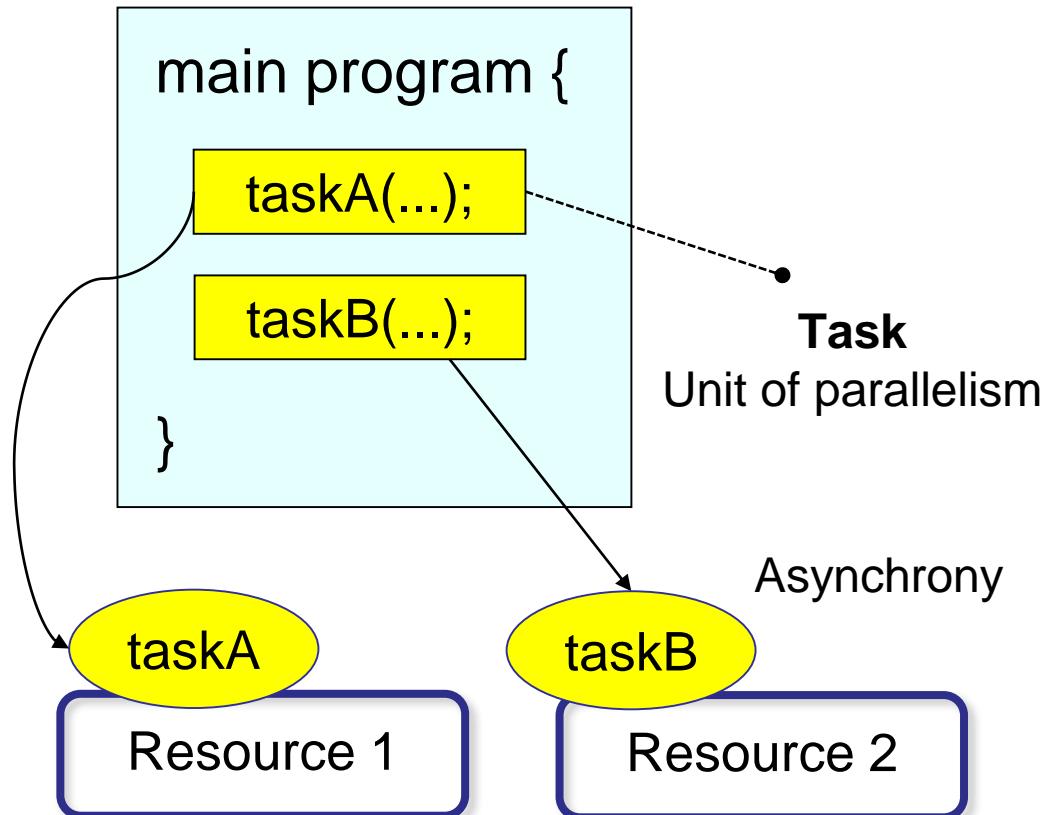
# Java Syntax

# COMPSs Bindings

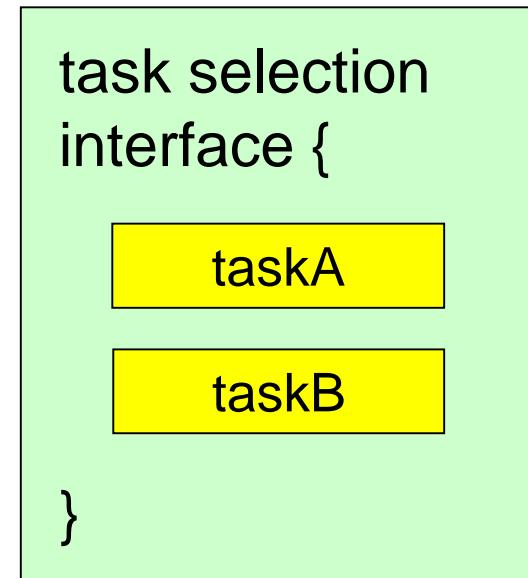


# Programming Model: Steps

## 1. Identify tasks



## 2. Select tasks

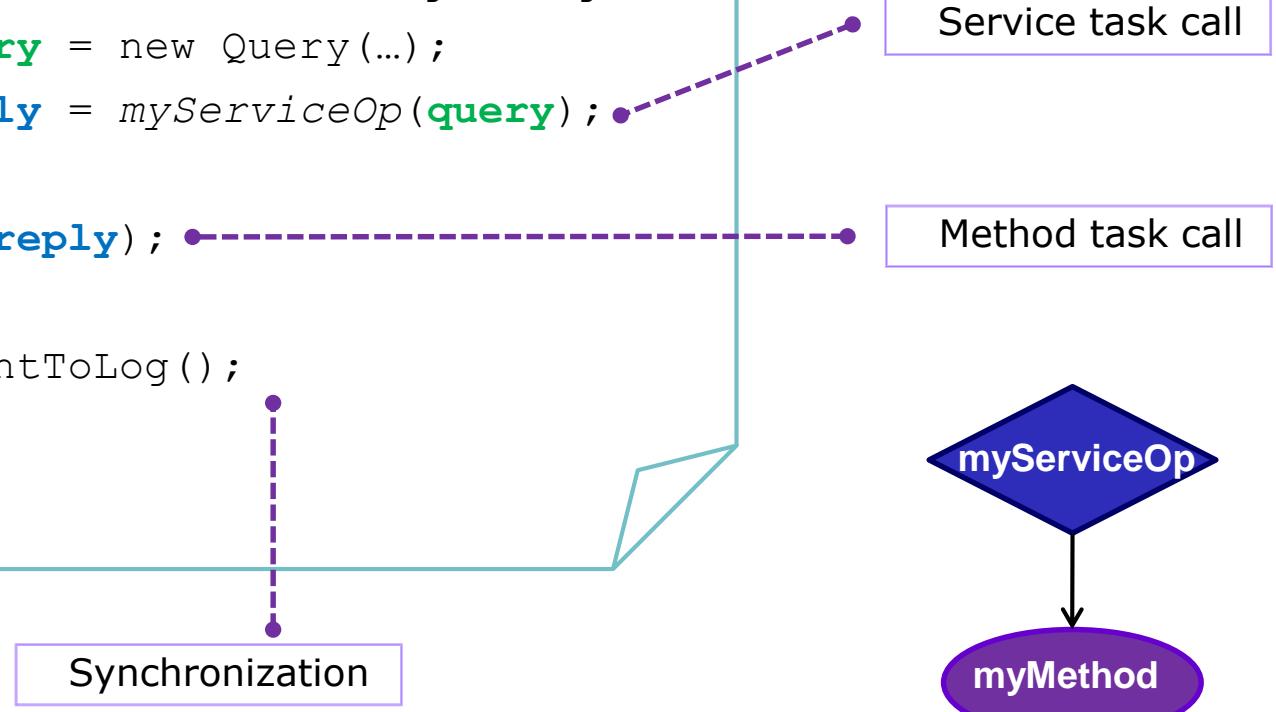


# Programming Model: Task selection interface

```
public interface SampleItf {  
    @Constraints(processorCPUCount = 1, memoryPhysicalSize = 0.5f)  
    @Method(declaringClass = "servicess.Example")  
    void myMethod(  
        @Parameter(direction = INOUT)  
        Reply r  
    );  
  
    @Service(namespace = "http://servicess.es/example",  
             name = "SampleService",  
             port = "SamplePort")  
    Reply myServiceOp(  
        @Parameter(direction = IN)  
        Query q  
    );  
}
```

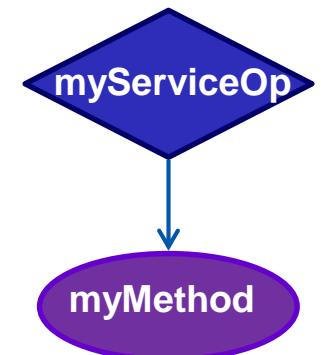
# Programming Model: Regular Main program

```
public class App {  
  
    public static void main(String[] args) {  
        Query query = new Query(...);  
        Reply reply = myServiceOp(query);  
  
        myMethod(reply);  
    }  
}
```



# Programming Model: Service Operation

```
public class ServiceApp {  
    @Orchestration  
    public static void sampleComposite() {  
        Query query = new Query(...);  
        Reply reply = myServiceOp(query);  
  
        myMethod(reply);  
  
        reply.printToLog();  
    }  
}
```



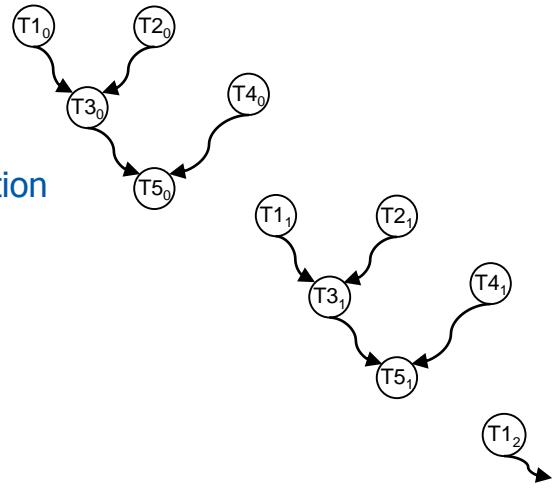
# Programming Model: Summary

Sequential Code

```
...  
for (i=0; i<N; i++) {  
    T1 (data1, data2);  
    T2 (data4, data5);  
    T3 (data2, data5, data6);  
    T4 (data7, data8);  
    T5 (data6, data8, data9);  
}  
...
```

(a) Task selection +  
parameters direction  
(input, output, inout)

(b) Task graph creation  
based on data  
dependencies



(d) Task completion,  
synchronization

(c) Scheduling,  
data transfer,  
task execution

Parallel Resources

Resource 1

Resource 2

...

Resource N





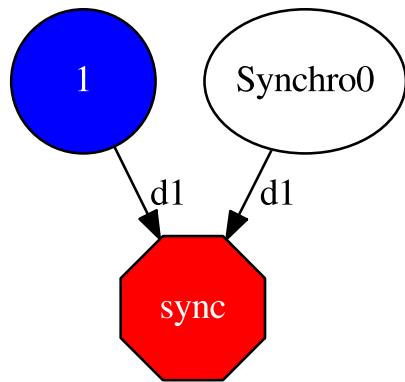
**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

# JAVA EXAMPLE

# Java COMPSs applications

## « Simple counter



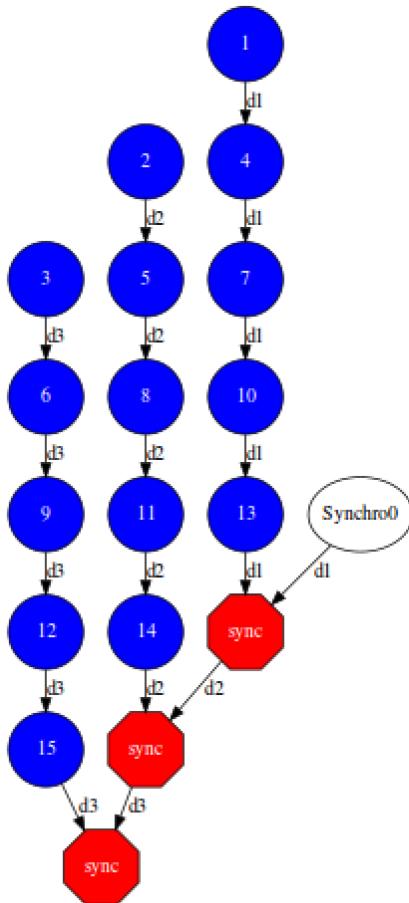
```
// Write value  
FileOutputStream fos = new FileOutputStream(fileName);  
    fos.write(initialValue);  
    fos.close();          System.out.println("Initial counter value  
is " + initialValue);  
  
//Execute increment  SimpleImpl.increment(fileName);
```

```
// Write new value  
FileInputStream fis = new FileInputStream(fileName);  
int finalValue = fis.read();  
fis.close();
```

```
public interface SimpleIf {  
    @Constraints(processorCoreCount = 1)  
    @Method(declaringClass = "simple.SimpleImpl")  
  
    void increment(  
        @Parameter(type = Type.FILE, direction = Direction.INOUT)  
        String file      );}
```

# Java COMPSs applications

## « Counter++



```
int N = Integer.parseInt(args[0]);
int counter1 = Integer.parseInt(args[1]);
int counter2 = Integer.parseInt(args[2]);
int counter3 = Integer.parseInt(args[3]);

// Initialize counter files
System.out.println("Initial counter values:");
initializeCounters(counter1, counter2, counter3);
// Print initial counters state
printCounterValues();

// Execute increment tasks
for (int i = 0; i < N; ++i) {
    IncrementImpl.increment(fileName1);
    IncrementImpl.increment(fileName2);
    IncrementImpl.increment(fileName3);
}

public static void increment(String counterFile) throws
FileNotFoundException, IOException {                                // Read value
    FileInputStream fis = new
    FileInputStream(counterFile);                                int count = fis.read();
    fis.close();                                                 // Write new
                                                              
value          FileOutputStream fos = new
FileOutputStream(counterFile);                                fos.write(++count);
fos.close(); }
```

# Matrix multiply in Java

```
for (int i = 0; i < MSIZE; i++) {  
    for (int j = 0; j < MSIZE; j++) {  
        for (int k = 0; k < MSIZE; k++) {  
            MatmulImpl.multiplyAccumulative( _C[i][j], _A[i][k], _B[k][j] );  
        }  
    }  
}
```

```
public static void multiplyAccumulative( String f3, String f1,  
String f2 )  
{  
    Block a = new Block( f1 );  
    Block b = new Block( f2 );  
    Block c = new Block( f3 );  
    c.multiplyAccum( a, b );  
    try  
    {...  
}  
  
public void multiplyAccum ( Block a, Block b )  
{  
    for( int i = 0; i < this.bRows; i++ )           // rows  
        for( int j = 0; j < this.bCols; j++ )       // cols  
            for ( int k = 0; k < this.bCols; k++ ) // cols  
                this.data[i][j] += a.data[i][k] * b.data[k][j];  
}
```

# Matrix multiply in Java

```
package matmul;

import integratedtoolkit.types.annotations.Constraints;
import integratedtoolkit.types.annotations.Method;
import integratedtoolkit.types.annotations.Parameter;
import integratedtoolkit.types.annotations.Parameter.*;

public interface MatmulItf {
    @Constraints(processorCoreCount = 4, memoryPhysicalSize = 1.5f)
    @Method(declaringClass = "matmul.MatmulImpl")
    void multiplyAccumulative(
        @Parameter(type = Type.FILE, direction = Direction.INOUT)
        String file1,
        @Parameter(type = Type.FILE, direction = Direction.IN)
        String file2,
        @Parameter(type = Type.FILE, direction = Direction.IN)
        String file3
    );
}
```



**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

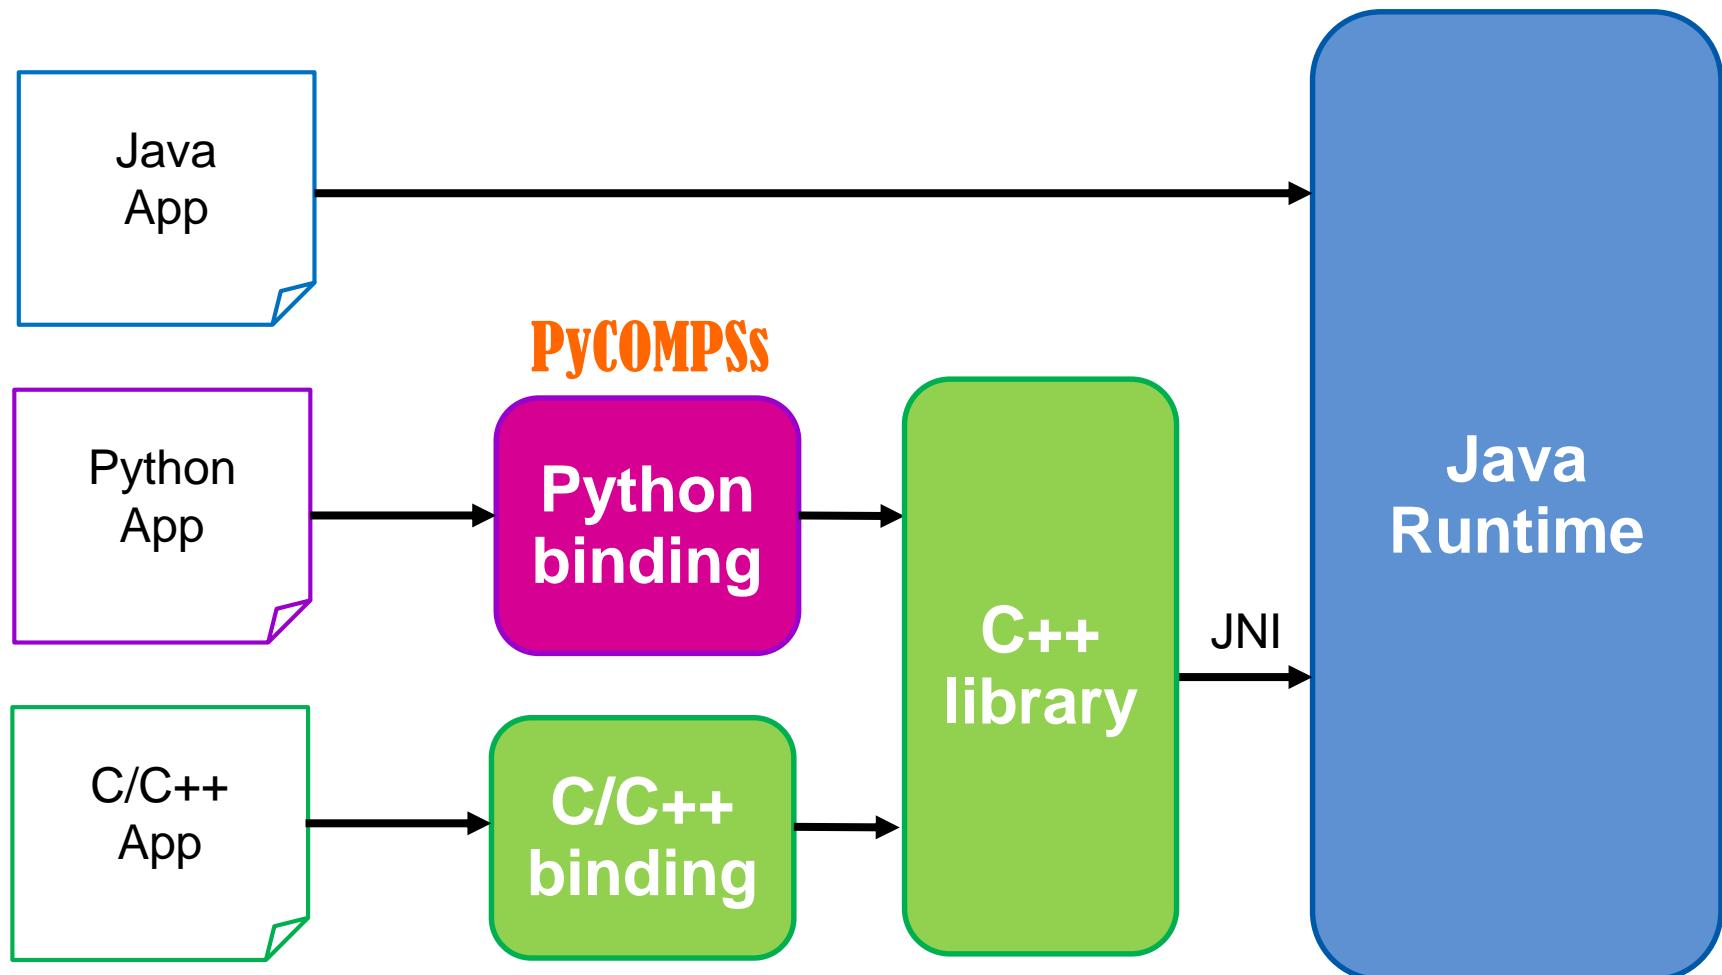
# Python Syntax

# Why Python?

- « Python is powerful... and fast;  
plays well with others;  
runs everywhere;  
is friendly & easy to learn;  
is Open. \*
- « Its design philosophy emphasizes code readability, and its syntax allows programmers to express concepts in fewer lines of code than would be possible in languages such as C
- « Large community using it, including scientific and numeric
- « Object-oriented programming and structured programming are fully supported
- « Large number of software modules available (38,000 as of January 2014)



# COMPSs Bindings



# PyCOMPSs: Task definition

## Task definition with Python decorators

- Provide information about task parameters (*TYPE\_DIRECTION*):
  - Type
    - Only mandatory for files
    - Inferred for the rest of the types
  - Direction
    - Default IN (read-only)
    - Mandatory for INOUT (read-write) and OUT (write-only)

The diagram illustrates the annotation of a Python function with task decorators. A dashed line with three dots connects the annotations to the corresponding parts of the code. The first dot points to the parameter `a` in the decorator `@task(a = INOUT, b = FILE_OUT)`, with the label "type inferred". The second dot points to the parameter `b` in the same decorator, with the label "explicit type and direction". The third dot points to the parameter `c` in the function definition `def my_func(a, b, c):`, with the label "type inferred, default direction (IN)".

```
@task(a = INOUT, b = FILE_OUT)
def my_func(a, b, c):
    ...
```

# PyCOMPSs: Task definition (II)

## « The @task decorator: special arguments

- Type of the return value → mandatory if a value is returned

```
@task(returns = int)
def ret_func():
    return 1
```

- Does the task modify the callee object? → default True

```
class MyClass(object):

    @task(isModifier = False)
    def instance_method(self):
        ... # self is NOT modified here
```

- Is it a priority task? → Default False

```
@task(priority = True)
def prio_func():

    ...
```

# PyCOMPSs: Task types

## « What can be selected as a task?

- (a) Functions
- (b) Instance methods
- (c) Class methods

```
@task( ... )
def my_function(...):
    ...
```

(a)

```
class Foo(object):

    @task( ... )
    def my_i_method(self, ...):
        ...

    @classmethod

    @task( ... )
    def my_c_method(cls, ...):
        ...
```

(b)

(c)

# PyCOMPSs: Main program → Synchronization API

- « Data created or updated by a task can be used in the main program of the application
  - But we need to synchronize first!
- « Two API methods for synchronization
  - `compss_open` → files

```
my_file = 'file.txt'  
func(my_file) •----- func is a task that modifies my_file  
fd = compss_open(my_file)  
  
...
```

- `compss_wait_on` → objects

```
my_obj = MyClass()  
my_obj.method() •----- method is a task that modifies my_obj  
my_obj = compss_wait_on(my_obj)  
  
...
```

# PyCOMPSs: Main program → Future objects

- « Mechanism to make asynchronous those tasks that return a value
  - Synchronization is only triggered when necessary
- « The future object is a representative of the object yet to be generated

```
@task(returns = MyClass)
def ret_func():
    return MyClass(...)
```

...

future object

```
o = ret_func()
```

# PyCOMPSs: Main program → Future objects (II)

- « A future object can be involved in a subsequent task call
  - PyCOMPSs will automatically enforce the dependency

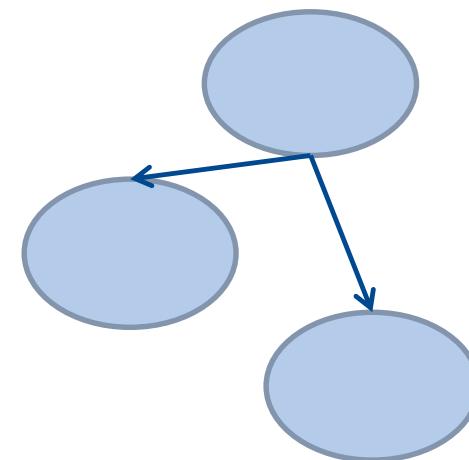
future object -----• `o = ret_func()`

...

`another_task(o)`

...

`o.yet_another_task()`



- « Synchronization from main program (same as other objects):

`o = ret_func()`

...

`o = compss_wait_on(o)`

# PyCOMPSs Constraints

- « Enables definition of tasks' constraints
  - Resource to execute the task should meet the constraint
- « Decorator definition:
  - `@constraint(constraint1="value1", constraint2="value2, ...)`
- « Examples of supported constraints:
  - ProcessorArch
  - ProcessorCoreCount
  - MemoryPhysicalSize
  - AppSoftware

# PyCOMPSs: Wrap-up example

« Invoke tasks as Python functions/methods

« API for data synchronization

« Task selection in function definition (decorators)

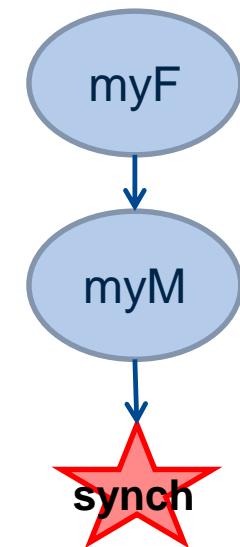
Main Program

```
foo = Foo()  
myFunction(foo)  
foo.myMethod()  
...  
foo = compss_wait_on(foo)  
foo.bar()
```

Function definition

```
@task(par = [INOUT])  
def myFunction(par):  
    ...
```

```
class Foo(object):  
    @task()  
    def myMethod(self):  
        ...
```



# Increment.py

Increment is an application that takes three different values and increases them a number of given times.

The purpose of this application is to show parallelism between the different increments.

Counters stored  
in files

```
def main_program():
    # Check and get parameters
    if len(sys.argv) != 5:
        usage()
        exit(-1)
    N = int(sys.argv[1])
    counter1 = int(sys.argv[2])
    counter2 = int(sys.argv[3])
    counter3 = int(sys.argv[4])

    # Initialize counter files
    initializeCounters(counter1, counter2, counter3)
    print "Initial counter values:"
    printCounterValues()

    # Execute increment
    for i in range(N):
        increment(FILENAME1)
        increment(FILENAME2)
        increment(FILENAME3)

    # Write final counters state (sync)
    print "Final counter values:"
    printCounterValues()
```

# Increment.py

Increment is an application that takes three different values and increases them a number of given times.

The purpose of this application is to show parallelism between the different increments.

```
def main_program():
    # Check and get parameters
    if len(sys.argv) != 5:
        usage()
        exit(-1)
    N = int(sys.argv[1])
    counter1 = int(sys.argv[2])
    counter2 = int(sys.argv[3])
    counter3 = int(sys.argv[4])

    # Initialize counter files
    initializeCounters(counter1, counter2, counter3)

    values:"
```

```
@task(filePath = FILE_INOUT)
```

```
def increment(filePath):
```

```
    # Read value
```

```
    fis = open(filePath, 'r')
```

```
    value = fis.read()
```

```
    fis.close()
```

```
    # Write value
```

```
    fos = open(filePath, 'w')
```

```
    fos.write(str(int(value) + 1))
```

```
    fos.close()
```

```
)
```

```
)
```

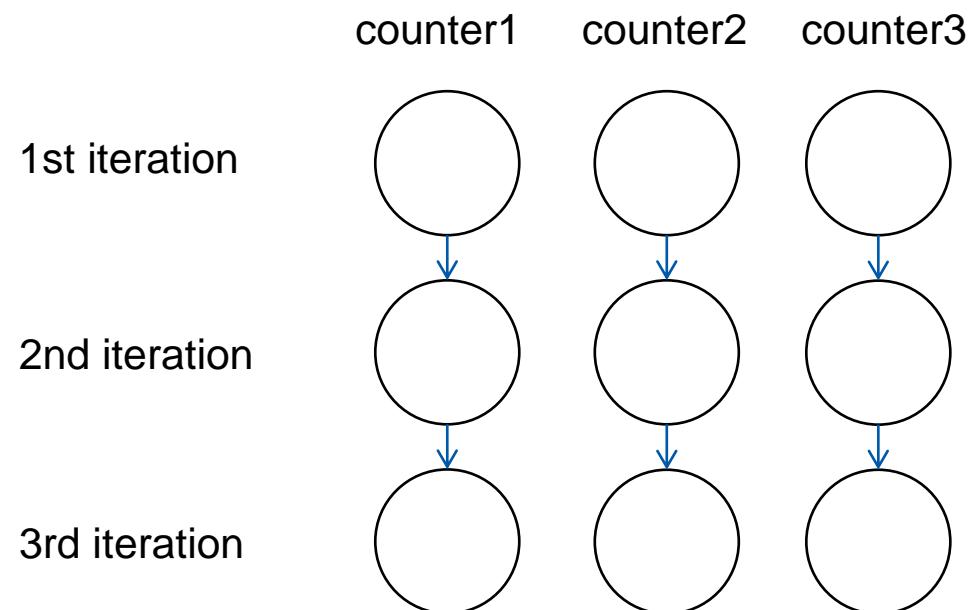
```
3)
```

```
state (sync)
```

```
values:"
```

## « Task graph identical to the Java case

### « Task Graph





**Barcelona  
Supercomputing  
Center**

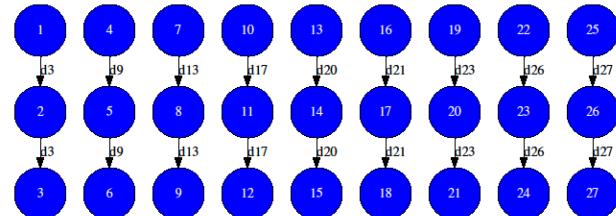
Centro Nacional de Supercomputación

## Other sample codes

# Matrix multiply with constraints in Python

```
from pycompss.api.constraint import constraint
from pycompss.api.task import task
from pycompss.api.parameter import INOUT

@constraint(ProcessorCoreCount=8)
@task(c = INOUT)
def multiply(a, b, c):
    import numpy
    c += a*b
```



```
args = sys.argv[1:]
MSIZE = int(args[0])
BSIZE = int(args[1])

A = B = C = []
# Initialize A, B & C as np.array(BSIZE, BSIZE)
initialize_variables()

for i in range(MSIZE):
    for j in range(MSIZE):
        for k in range(MSIZE):
            multiply(A[i][k], B[k][j], C[i][j])
```

# PyCOMPSs constraints

```
@task(returns=str)
def mutate(source):
    charpos = random.randint(0, len(source) - 1)
    parts = list(source)
    parts[charpos] = chr(ord(parts[charpos]) +
                         random.randint(-1,1))
    return ''.join(parts)
```

```
@constraint(AppSoftware="Numpy")
@task(returns=int)
def fitness(source, target):
    import numpy
    fitval = 0
    a = numpy.array([ord(i) for i in source])
    b = numpy.array([ord(i) for i in target])
    fitval = numpy.linalg.norm(a-b)
    return fitval
```

```
fitval = fitness(source, target)
i = 0
while True:
    i += 1
    m = mutate(source)
    fitval_m = fitness(m, target)
    fitval_m = compss_wait_on(fitval_m)
    if fitval_m < fitval:
        fitval = fitval_m
    m = compss_wait_on(source)
```

# Sample code: Kmeans @ PyCOMPSS

```
from pycompss.api.api import compss_wait_on
size = int(numV / numFrag)

X = [genFragment(size, dim) for _ in range(numFrag)]
mu = init_random(dim, k)
oldmu = []
n = 0
startTime = time.time()
while not has_converged(mu, oldmu, epsilon, n, maxIterations):
    oldmu = mu
    clusters = [cluster_points_partial(X[f], mu, f * size) for f in range(numFrag)]
    partialResult = [partial_sum(X[f], clusters[f], f * size) for f in range(numFrag)]

    mu = merge_reduce(reduceCentersTask, partialResult)
    mu = compss_wait_on(mu)
    mu = [mu[c][1] / mu[c][0] for c in mu]
    n += 1
return (n, mu)
```

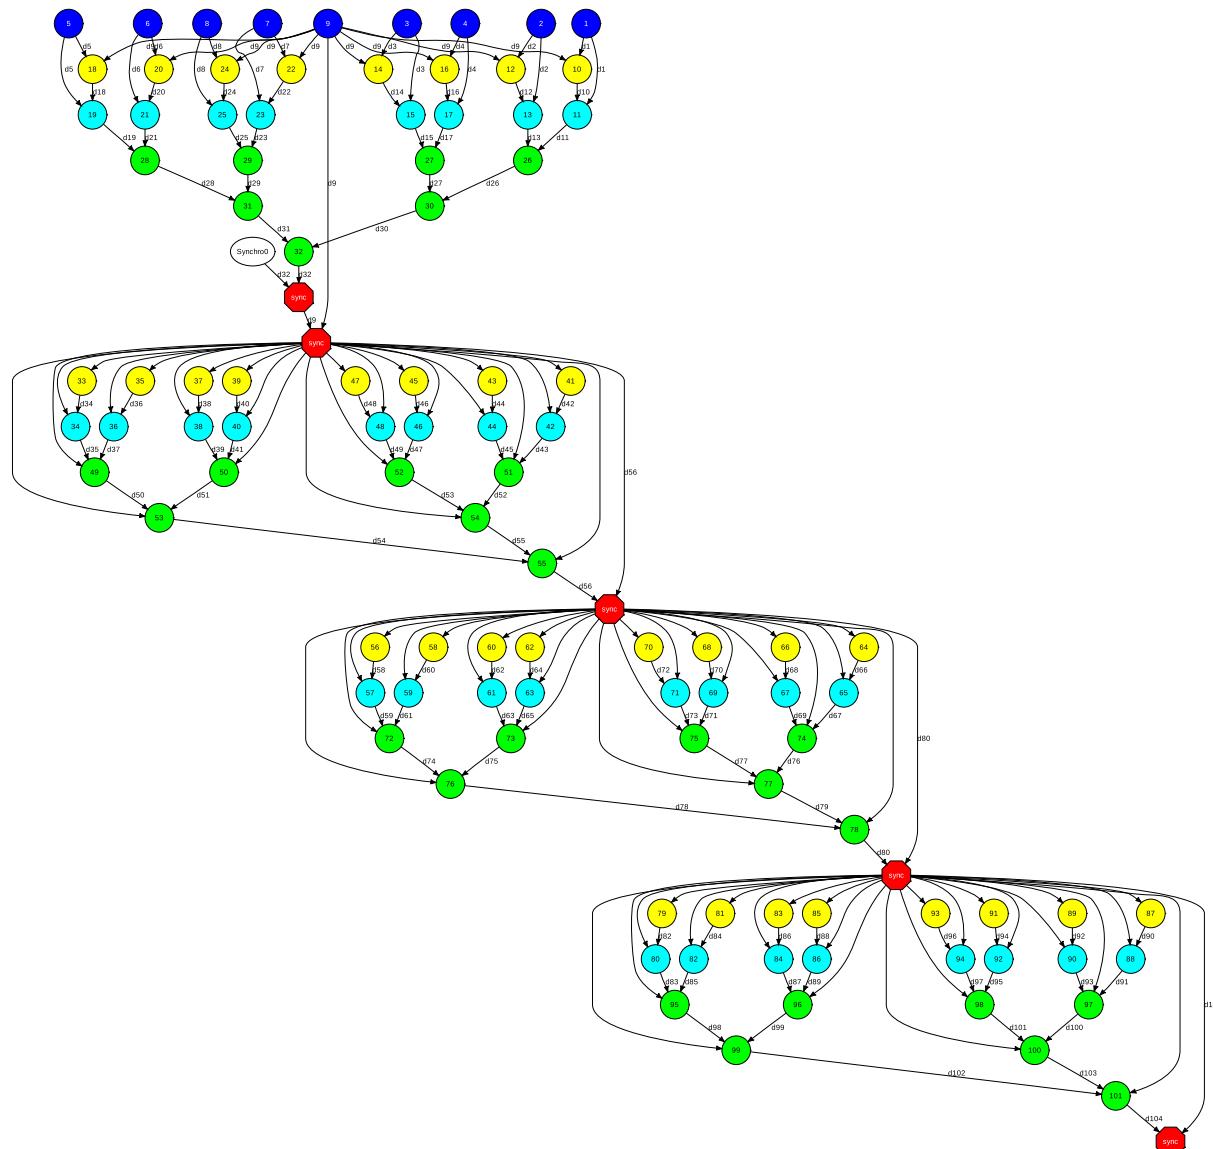
```
@task(returns=dict)
def partial_sum(XP, clusters, ind):
    import numpy as np
    XP = np.array(XP)
    p = [(i, [(XP[j] - ind) for j in clusters[i]]) for i in clusters]
    dic = {}
    for i, l in p:
        dic[i] = (len(l), np.sum(l, axis=0))
    return dic
```

```
@task(returns=dict, priority=True)
def reduceCentersTask(a, b):
    for key in b:
        if key not in a:
            a[key] = b[key]
        else:
            a[key] = (a[key][0] + b[key][0],
                      a[key][1] + b[key][1])
    return a
```

```
@task(returns=dict)
def cluster_points_partial(XP, mu, ind):
    import numpy as np
    dic = {}
    XP = np.array(XP)
    for x in enumerate(XP):
        bestmukey = min([(i[0], np.linalg.norm(x[1] - mu[i[0]]))
                        for i in enumerate(mu)], key=lambda t: t[1])[0]
        if bestmukey not in dic:
            dic[bestmukey] = [x[0] + ind]
        else:
            dic[bestmukey].append(x[0] + ind)
    return dic
```

# Sample code: Kmeans @ PyCOMPSS

- Task graph:
  - 8 fragments
  - 4 iterations



- Computation of mutual cross-correlations between all pairs of a set of spike data
- Also computes the cross-correlations for surrogate data sets for each neuron pair



```

f = open('./spikes.dat', 'r')
spikes = pickle.load(f)
f.close()
#preallocate result variables
num_ccs = (num_neurons**2 - num_neurons)/2
cc_orig = zeros((num_ccs,2*maxlag+1))
cc_surrs = zeros((num_ccs,2*maxlag+1,num_surrs))
idxrange = range(num_bins-maxlag,num_bins+maxlag+1)
row = 0

#for all pairs ni,nj such that nj > ni
for ni in range(num_neurons-1):
    for nj in range(ni+1,num_neurons):
        cc_orig[row,:] = correlate(spikes[ni,:],spikes[nj,:],...)
        num_spikes_i = sum(spikes[ni,:])
        num_spikes_j = sum(spikes[nj,:])
        for surrogate in range(num_surrs):
            surr_i = zeros(num_bins)
            surr_i[random.random_integers(0,num_bins-1,num_spikes_i)] = 1
            surr_j = zeros(num_bins)
            surr_j[random.random_integers(0,num_bins-1,num_spikes_j)] = 1
            cc_surrs[row,:,:surrogate] = correlate(surr_i,surr_j,"full")[idxrange]
        row = row +1

#save results
f = open('./result_cc_originals.dat','w')
pickle.dump(cc_orig,f)
f.close()
f = open('./result_cc_surrogates.dat','w')
pickle.dump(cc_surrs,f)
f.close()

```

Sequential code

# Neuroscience Data Processing @ Parallel Python

## Main program

```
...  
# tuple of all parallel python servers to connect with  
ppservers = ('comp1.my-network', 'comp2.my-network' ...  
  
if len(sys.argv) > 1:  
    ncpus = int(sys.argv[1])  
    #creates jobserver with ncpus workers  
    job_server = pp.Server(ncpus, ...  
else:  
    #creates jobserver with workers automatically detected  
    job_server = pp.Server(ppservers=ppservers, ...  
  
#wait for servers to come up  
time.sleep(5)  
  
#calculate number of nodes in total  
nlocalworkers = job_server.get_ncpus()  
activenodes = job_server.get_active_nodes()  
workerids = activenodes.keys()  
nworkers = sum([activenodes[workerids[i]] for i in  
range(len(workerids))]) + nlocalworkers  
num_ccs = (num_neurons**2 - num_neurons)/2  
  
#calculate number of pairs each worker should process  
step = ceil(float(num_ccs)/nworkers)  
start_idx = 0  
end_idx = 0  
starts = zeros((nworkers+1,))  
...
```

## Explicit resources declaration

```
def cc_surrogate_range(start_idx, end_idx, seed, num_neurons,  
num_surrs, num_bins, maxlag):  
    ...
```

## Function definition

## Main program (cont)

```
for worker in range(nworkers):  
    start_idx = end_idx  
    end_idx = int(min((worker+1)*step,num_ccs))  
  
    ...  
    depfuncs = ()  
    depmodules = "numpy","pickle",  
    jobs.append(job_server.submit(cc_surrogate_range,...
```

## Explicit fork join

```
cc_original = zeros((num_ccs,2*maxlag+1))  
cc_surrs = zeros((num_ccs,2*maxlag+1,2))  
for worker in arange(nworkers):  
    start = starts[worker]  
    end = starts[worker + 1]  
    result = jobs[worker]()  
    cc_original[start:end,:] = result[0]  
    cc_surrs[start:end,:,:] = result[1]
```

## Data back

```
f = open('./result_cc_origins.dat','w')  
pickle.dump(cc_original,f)  
f.close()  
f = open('./result_cc_surrogates_conf.dat','w')  
pickle.dump(cc_surrs,f)  
f.close()
```

# Neuroscience Data Processing @ PyCOMPSS

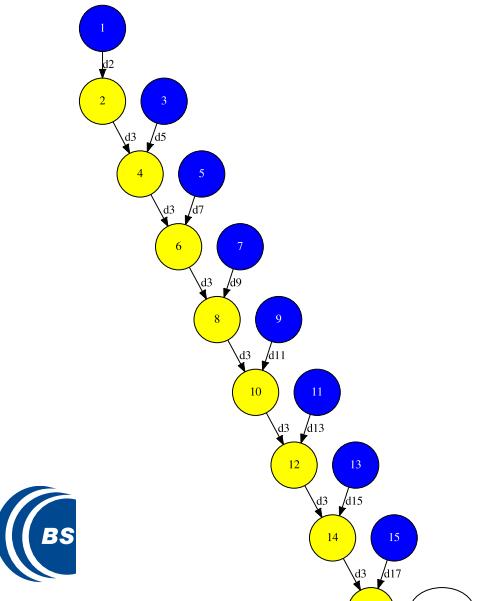
```
Main program
```

```
Import sys
from pycompss.api.api import compss_wait_on

num_frags = int(sys.argv[1])

#calculate number of pairs per fragment
num_ccs = (num_neurons**2 - num_neurons)/2
step = ceil(float(num_ccs)/num_frags)
start_idx = 0
end_idx = 0

seed = 2398645
delta = 1782324
```



```
@task(returns = list)
def cc_surrogate_range(start_idx, end_idx, seed, num_neurons,
                      num_surrs, num_bins, maxlag):
    ...
```

```
Main program (cont)
```

```
cc_original = zeros((num_ccs,2*maxlag+1))
cc_surrs = zeros((num_ccs,2*maxlag+1,2))
for frag in range(num_frags):
    start_idx = end_idx
    end_idx = int(min((frag+1)*step,num_ccs))
    result = cc_surrogate_range(start_idx, end_idx, seed, ...
                                num_neurons, num_surrs, num_bins, maxlag)
    gather(result, cc_original, cc_surrs, start_idx, end_idx)
    seed = seed + delta
```

```
f = open('./result_cc_origins.dat','w')
cc_original = compss_wait_on(cc_original)
pickle.dump(cc_original,f)
f.close()
```

```
f = open('./result_cc_surrogates_conf.dat','w')
cc_surrs = compss_wait_on(cc_surrs)
pickle.dump(cc_surrs,f)
f.close()
```



**Barcelona  
Supercomputing  
Center**

*Centro Nacional de Supercomputación*

# COMPSs execution environment examples

Application

Task Selection Interface



## How can I select the execution platform?



Grid

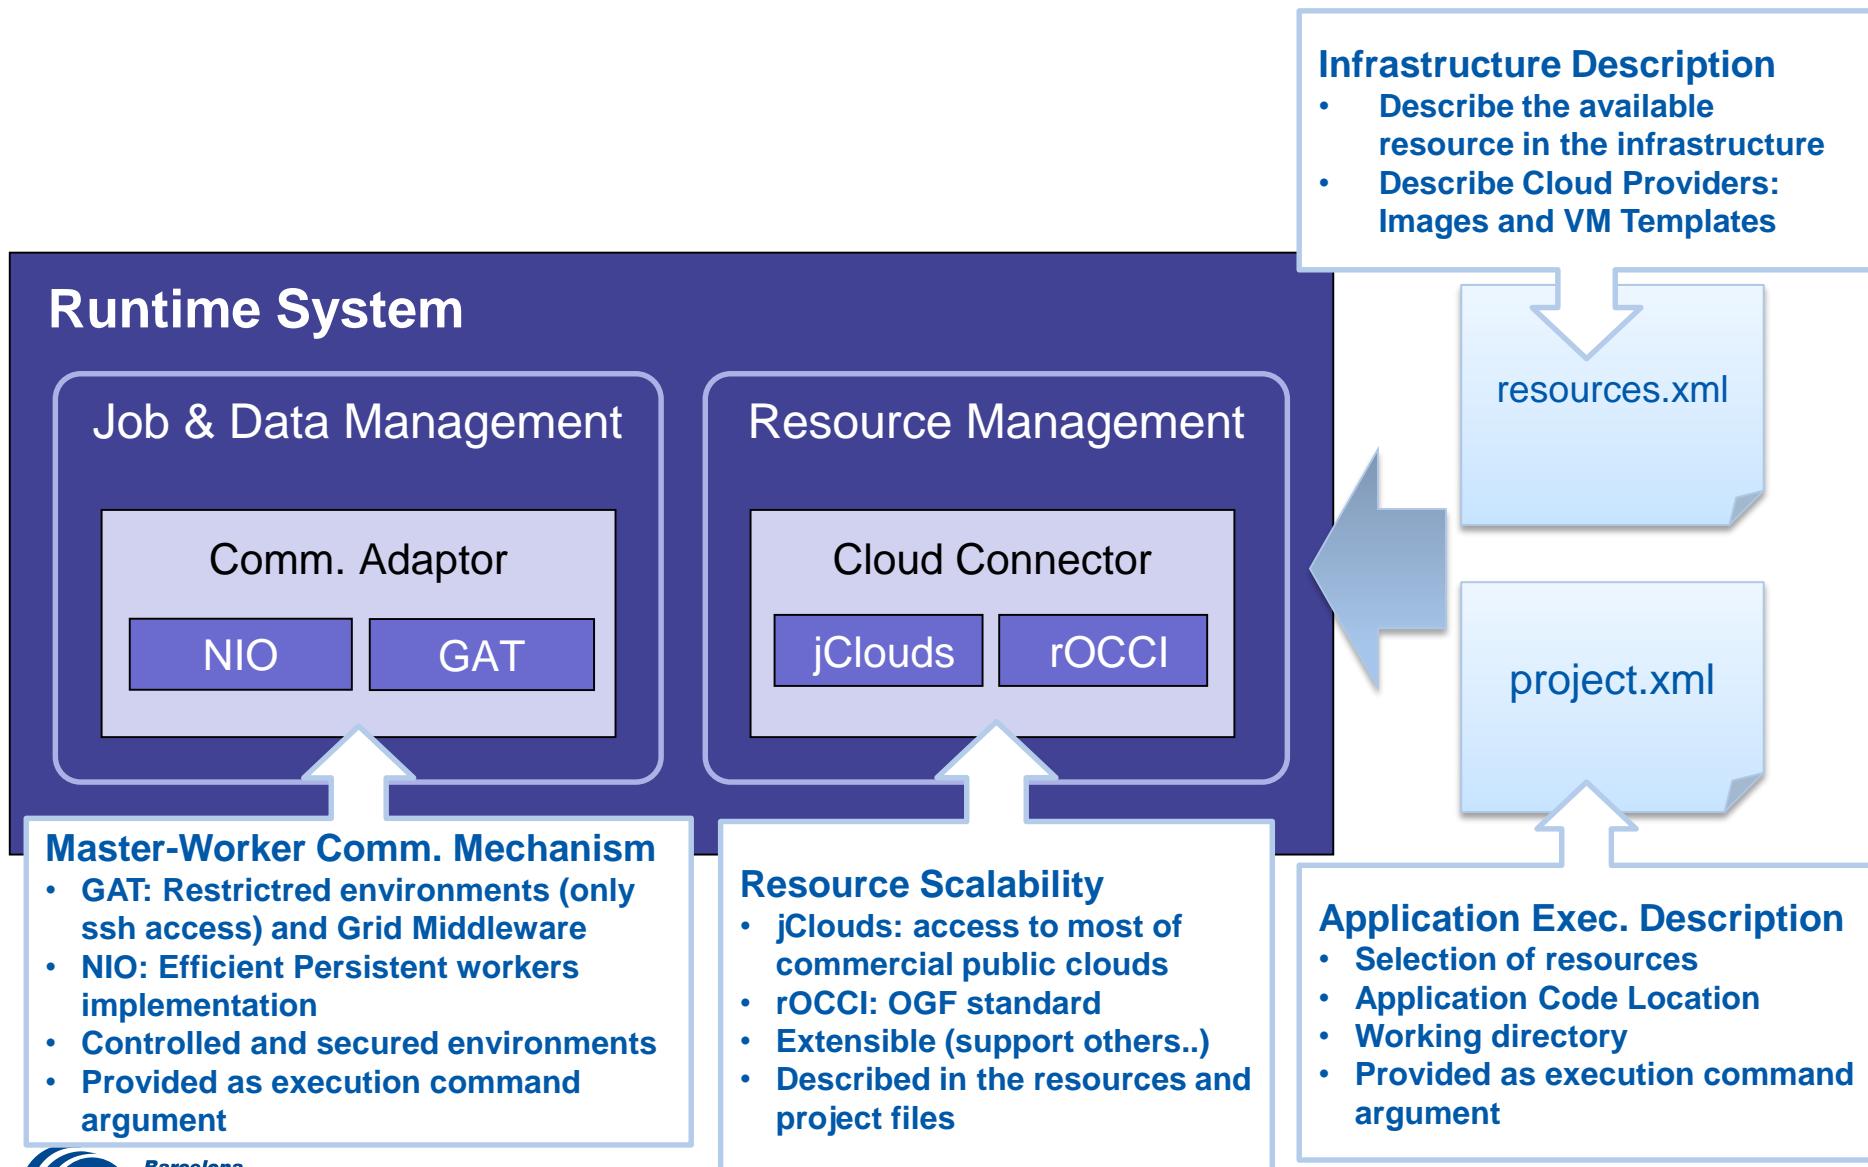


Cluster



Cloud

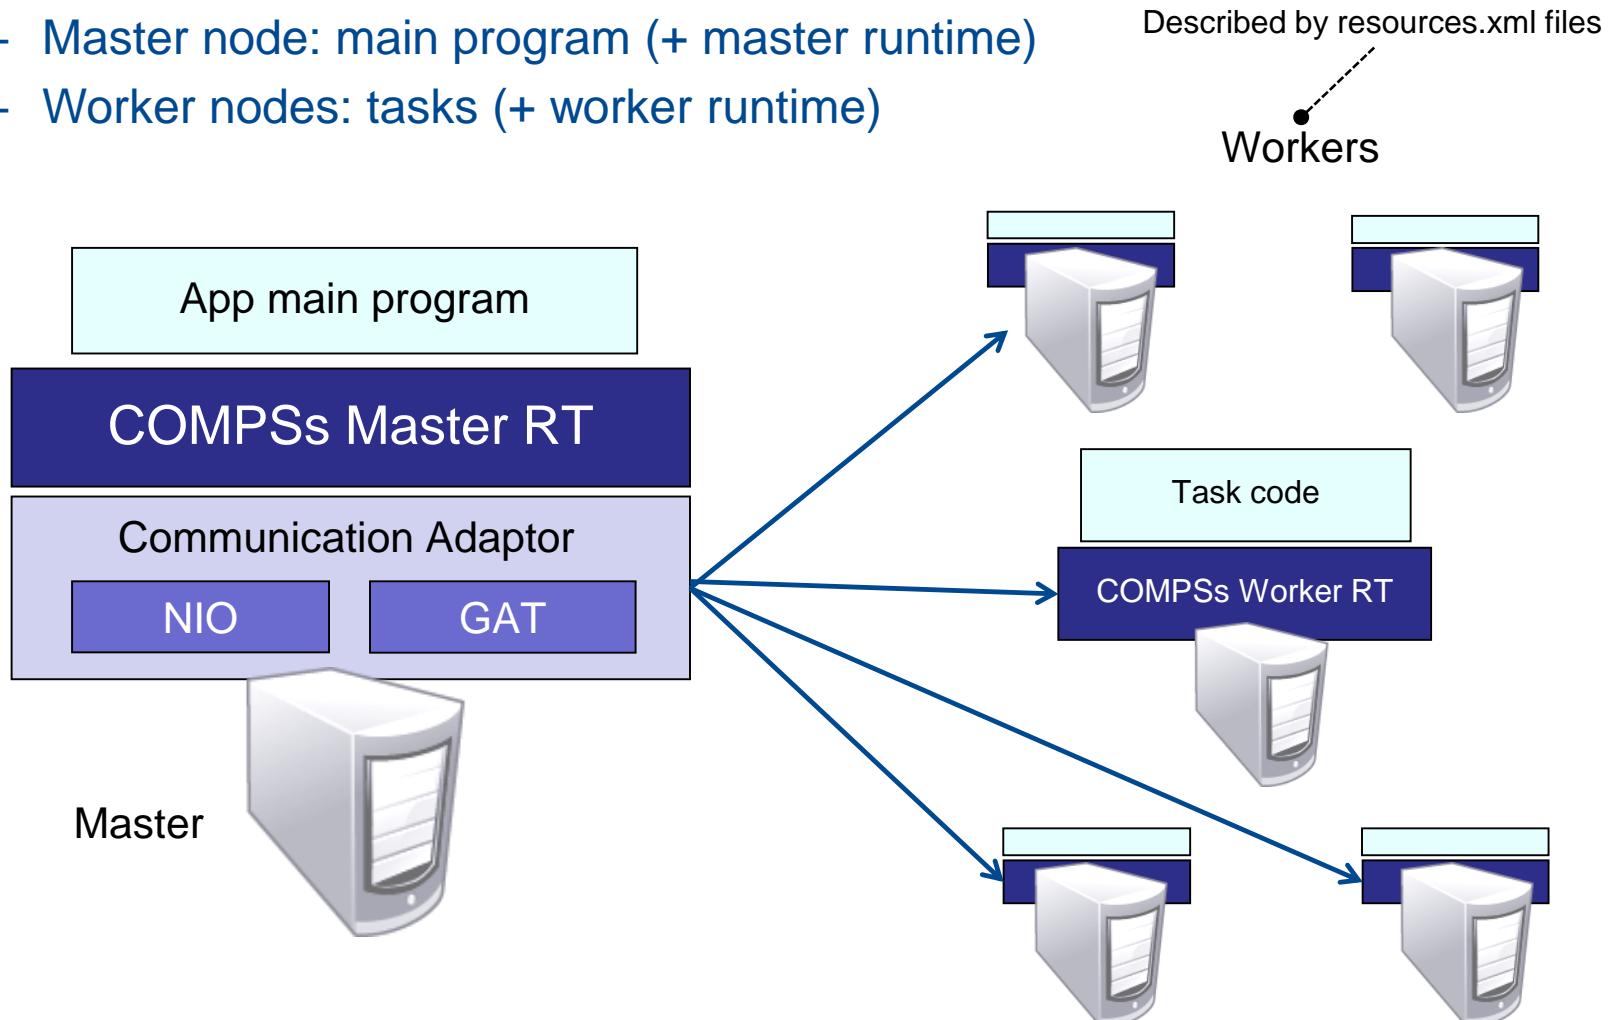
# Execution Environment Configuration Overview



# COMPSs in remote hosts (interactive)

## Typical setup:

- Master node: main program (+ master runtime)
- Worker nodes: tasks (+ worker runtime)



# Configuration: Resources Specification

## Resources.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<ResourceList>
    <!--Description for any physical node-->
    <Resource Name="172.20.200.18">
        <Capabilities>
            <Host>
                <TaskCount>0</TaskCount>
                <Queue>short</Queue>
                <Queue/>
            </Host>
            <Processor>
                <Architecture>x86</Architecture>
                <Speed>3.0</Speed>
                <CoreCount>2</CoreCount>
            </Processor>
            <OS>
                <OSType>Linux</OSType>
                <MaxProcessesPerUser>32</MaxProcessesPerUser>
            </OS>
            <StorageElement>
                <Size>30</Size>
            </StorageElement>
        ...
    </Resource>
```

```
    ...
    <Memory>
        <PhysicalSize>1</PhysicalSize>
        <VirtualSize>8</VirtualSize>
    </Memory>
    <ApplicationSoftware>
        <Software>Java</Software>
    </ApplicationSoftware>
    <Service/>
    <VO/>
    <Cluster/>
    <FileSystem/>
    <NetworkAdaptor/>
    <JobPolicy/>
    <AccessControlPolicy/>
    </Capabilities>
    <Requirements/>
</Resource>
<Resource Name="172.20.200.19">
    ...
</Resource>
<ResourceList>
```

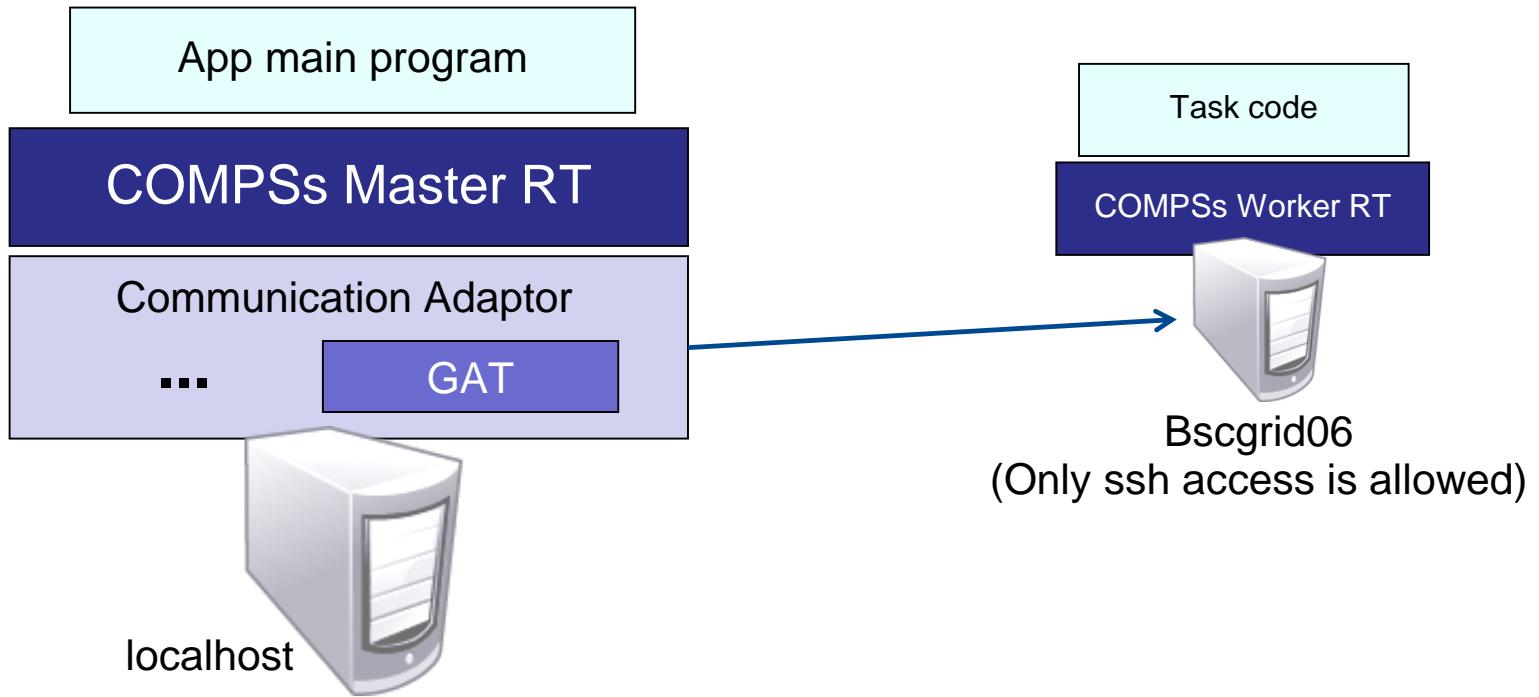
# Configuration: Project Specification

## Project.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<Project>
    <!--Description for any physical node-->
    <Worker Name="172.20.200.18">
        <InstallDir>/opt/COMPSS/Runtime/scripts/</InstallDir>
        <WorkingDir>/tmp/</WorkingDir>
        <User>user</User>
        <LimitOfTasks>1</LimitOfTasks>
    </Worker>

    <Worker Name="172.20.200.19">
        ...
    </Worker>
    ...
</Project>
```

# Example: COMPSs using Remote hosts (interactive)



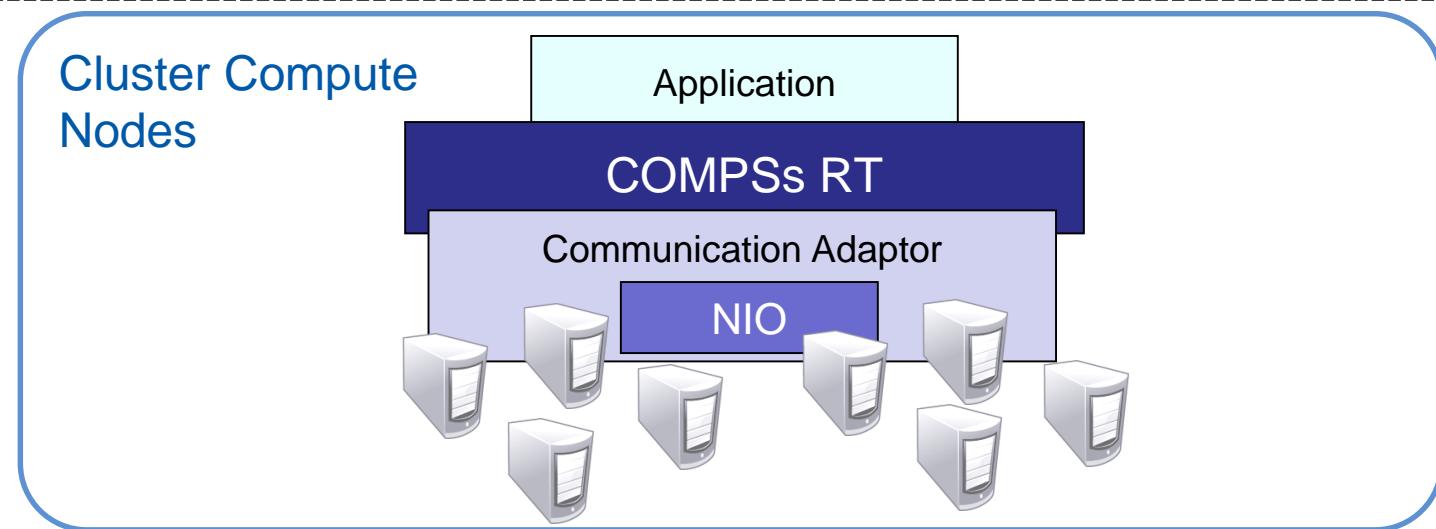
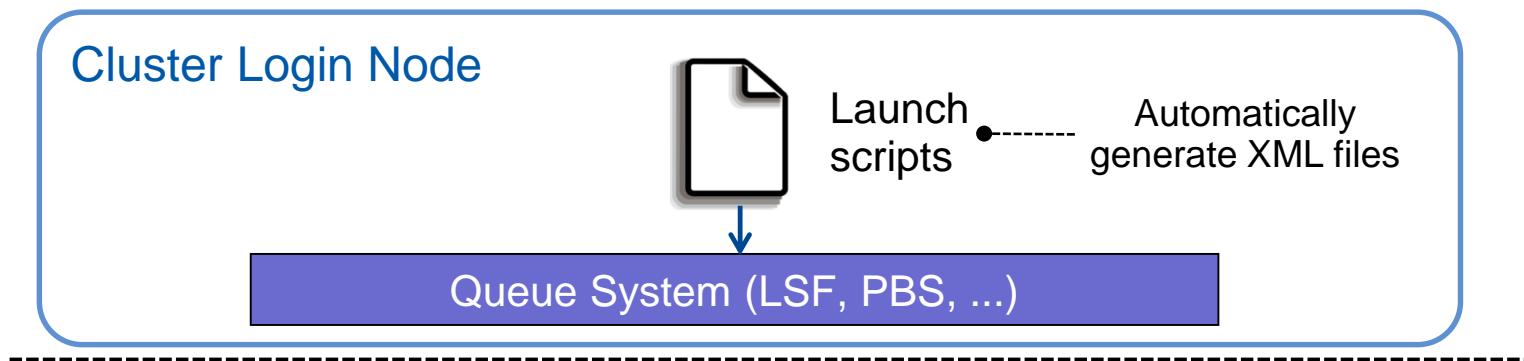
## « Steps:

- Deploy code in worker
- Run the application with specific resources and project.xml and GAT the adaptor

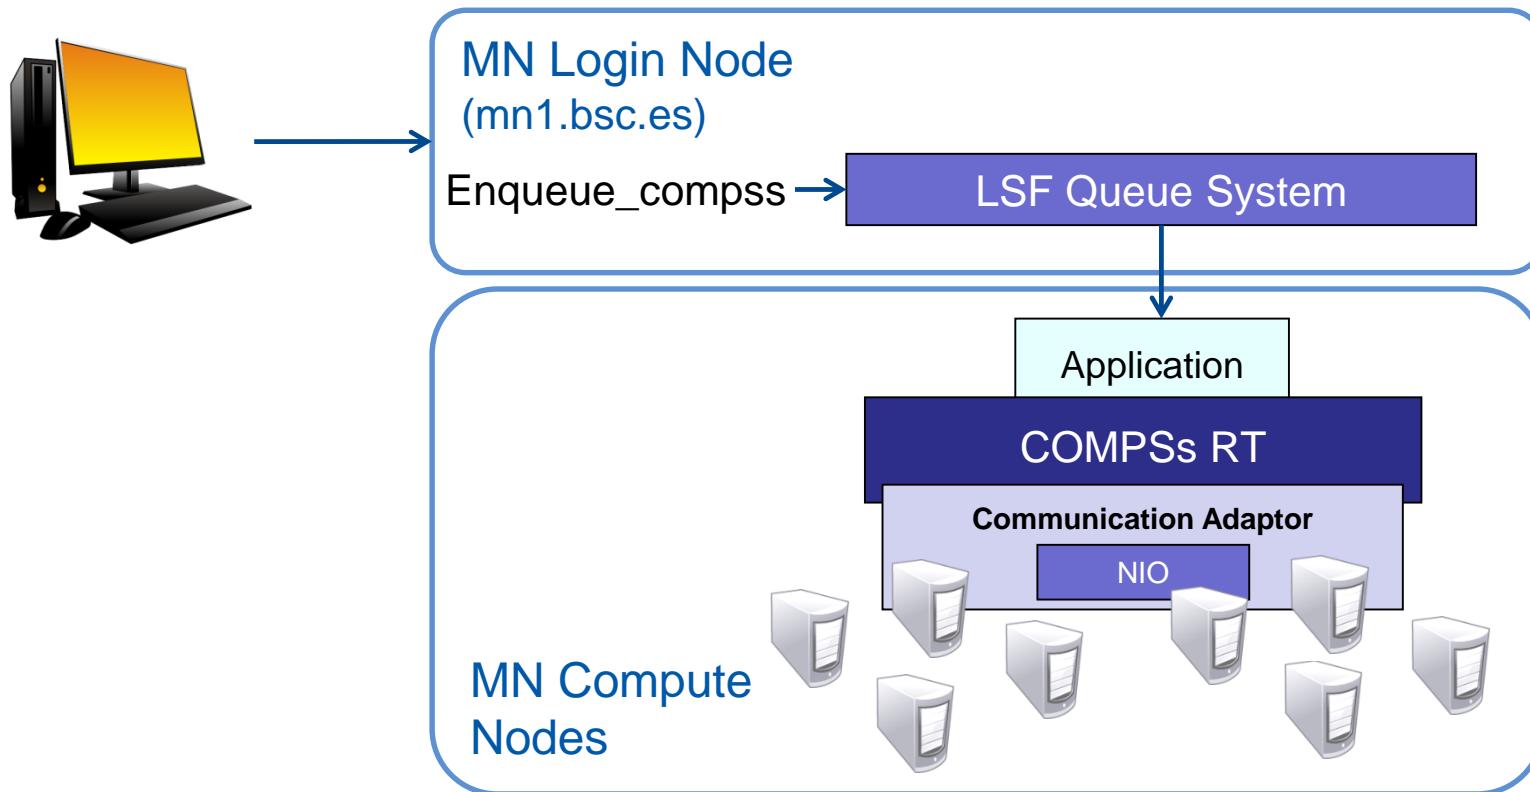
# COMPSs in a Cluster (queue system)

## Execution divided in two phases

- Launch scripts queue a whole COMPSs app execution
- Actual execution starts when reservation is obtained



# COMPSs in a Cluster (queue system)



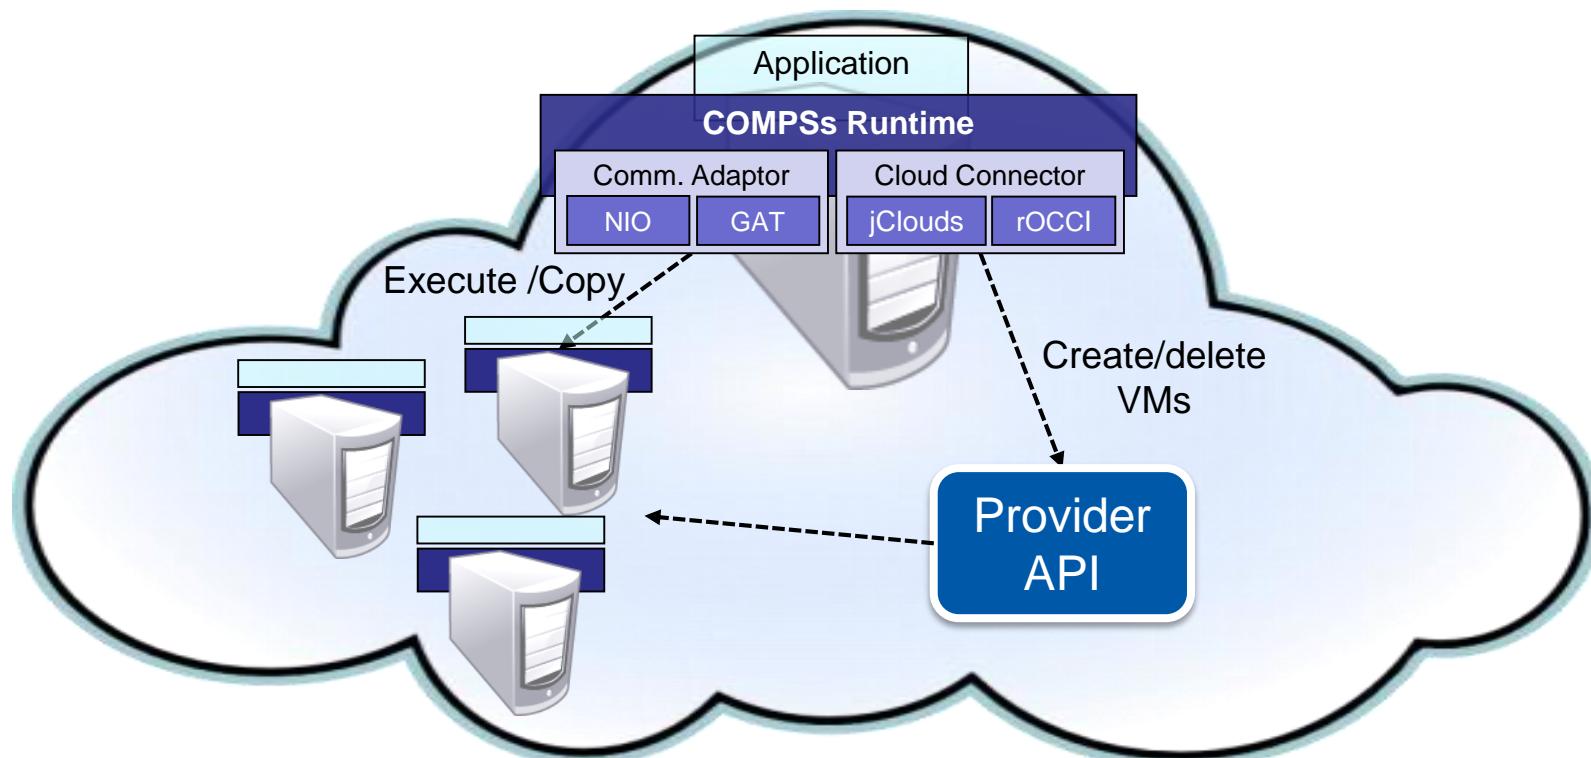
## Steps:

- Deploy app in MN
- Connect login node
- Launch “enqueue\_comps” script requesting the number of nodes

# COMPSs in Clouds

## Execution of COMPSs applications in Clouds

- Select the connector to interact with the Cloud provider
- Adaptor to communicate VMs (NIO if provider supports firewall management, GAT if only ssh)



# Cloud Configuration: Resources Specification

## Resources.xml

```
<ResourceList>
  <CloudProvider name="BSCCloud">
    <Server>https://bscgrid20.bsc.es:11443</Server>
    <Connector>integratedtoolkit.connectors.rocci.ROCCI</Connector>
    <ImageList>
      <Image name="debianbase"/>
    </ImageList>
    <CreationTime>120</CreationTime>
    <InstanceTypes>
      <Resource Name="bsc.small">
        <Capabilities>
          <Processor>
            <CoreCount>1</CoreCount>
          </Processor>
          <StorageElement>
            <Size>10.0</Size><!-- GB -->
          </StorageElement>
          <Memory>
            <PhysicalSize>1</PhysicalSize>><!-- GB -->
          </Memory>
        </Capabilities>
      </Resource>
      <Resource Name="bsc.medium">
        ...
      </Resource>
    </InstanceTypes>
  </CloudProvider>
</ResourceList>
```

# Cloud Configuration: Project Specification

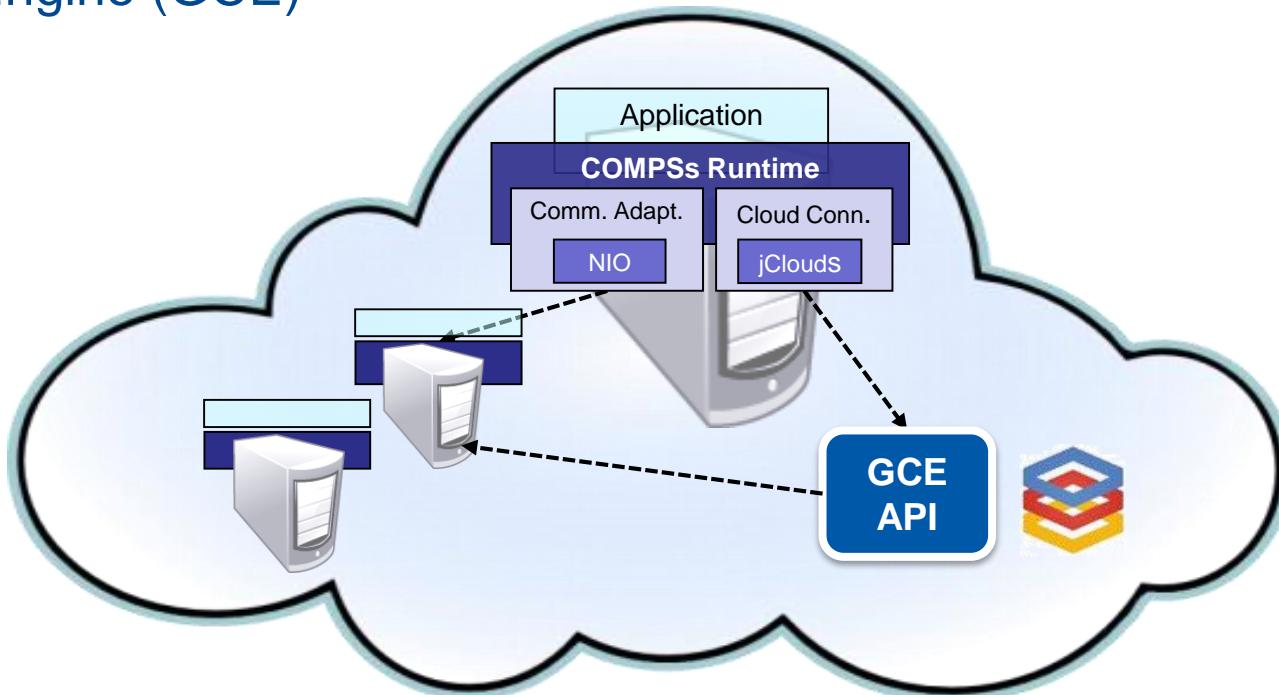
## Project.xml

```
<Project>
  <Cloud>
    <InitialVMs>0</InitialVMs>
    <minVMCount>2</min VMCount>
    <maxVMCount>5</max VMCount>
    <Provider name="BSCCloud">
      <LimitOfVMs>5</LimitOfVMs>
      <Property>
        <Name>user-cred</Name>
        <Value>/home/.../cert.pem</Value>
      </Property>
      <Property>
        <Name>user</Name>
        <Value>userbsc</Value>
      </Property>
      ...
    ...
```

```
      ...
      <ImageList>
        <Image name="debianbase">
          <InstallDir>/opt/COMPSS/Runtime/scripts</InstallDir>
          <WorkingDir>/tmp/</WorkingDir>
          <User>user</User>
          <Package>
            <Source>/home/.../AppName.tar.gz</Source>
            <Target>/home/user</Target>
          </Package>
        </Image>
      </ImageList>
      ...
      <InstanceTypes>
        <Resource name="bsc.small"/>
      </InstanceTypes>
    </Provider>
  </Cloud>
</Project>
```

# DEMO: COMPSs in Clouds

- Execution of COMPSs applications in Google Compute Engine (GCE)



[https://www.youtube.com/watch?v=XGaqUje\\_2zY](https://www.youtube.com/watch?v=XGaqUje_2zY)

## « Combine Different Environment

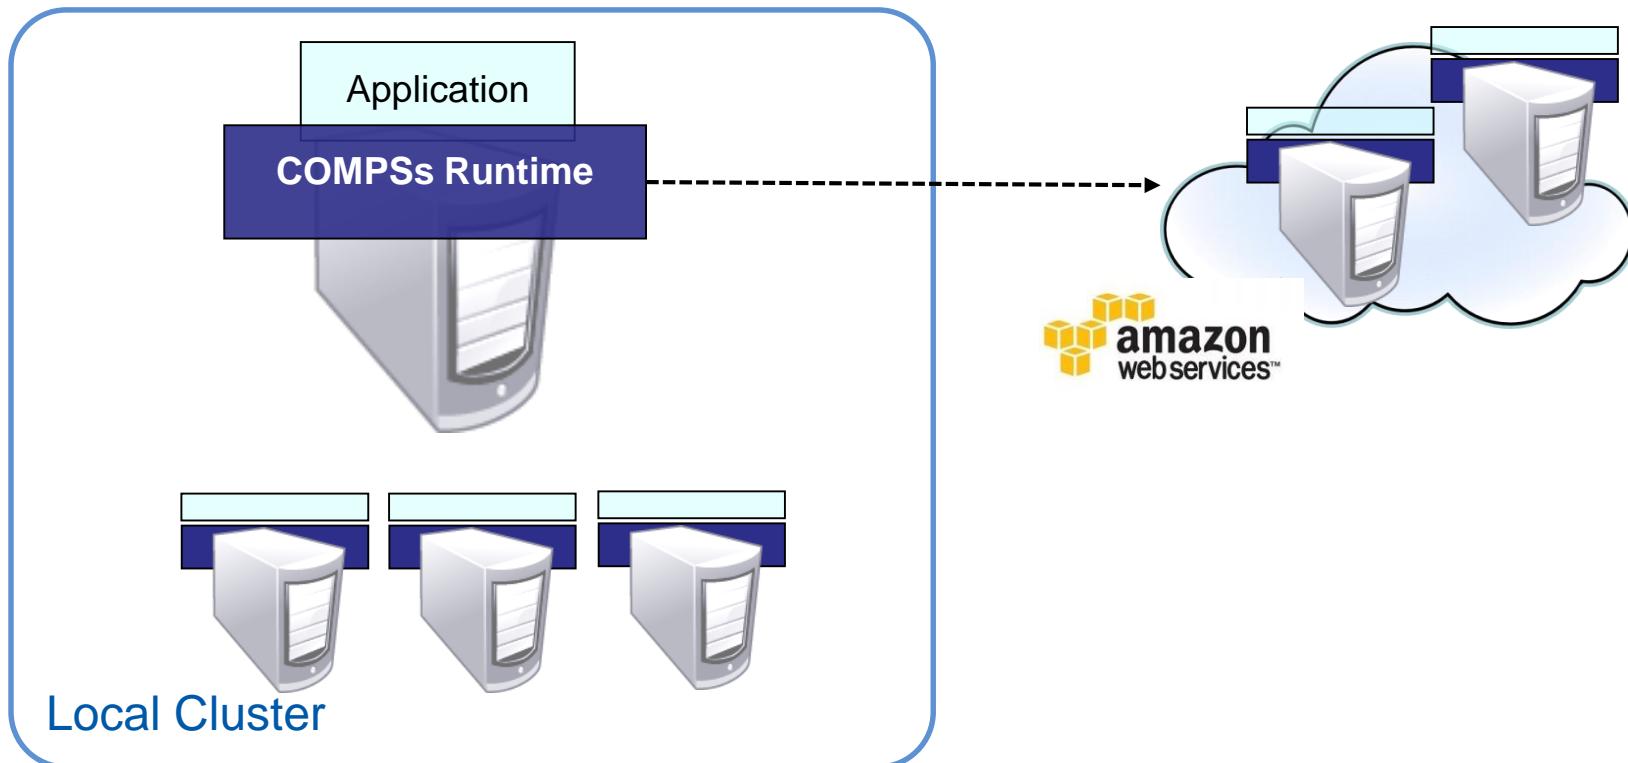
- Cloud Bursting
- Multiple Grids

## « COMPSs for scaling Web Service

## « Real Applications implemented with COMPSs

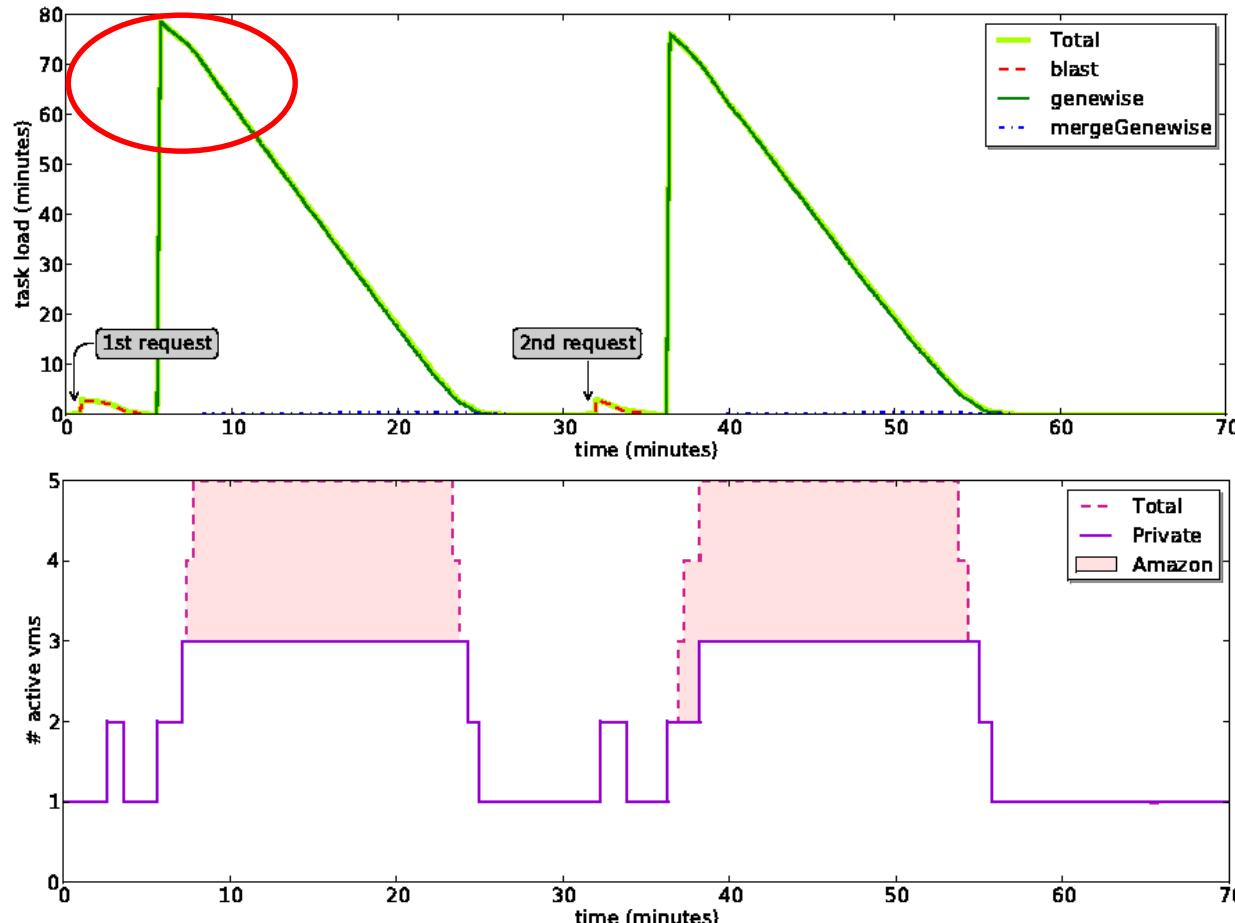
# Cloud Bursting

- Execution of COMPSS applications in Clouds
  - Select de connector to interact Cloud providers connectors, SSH

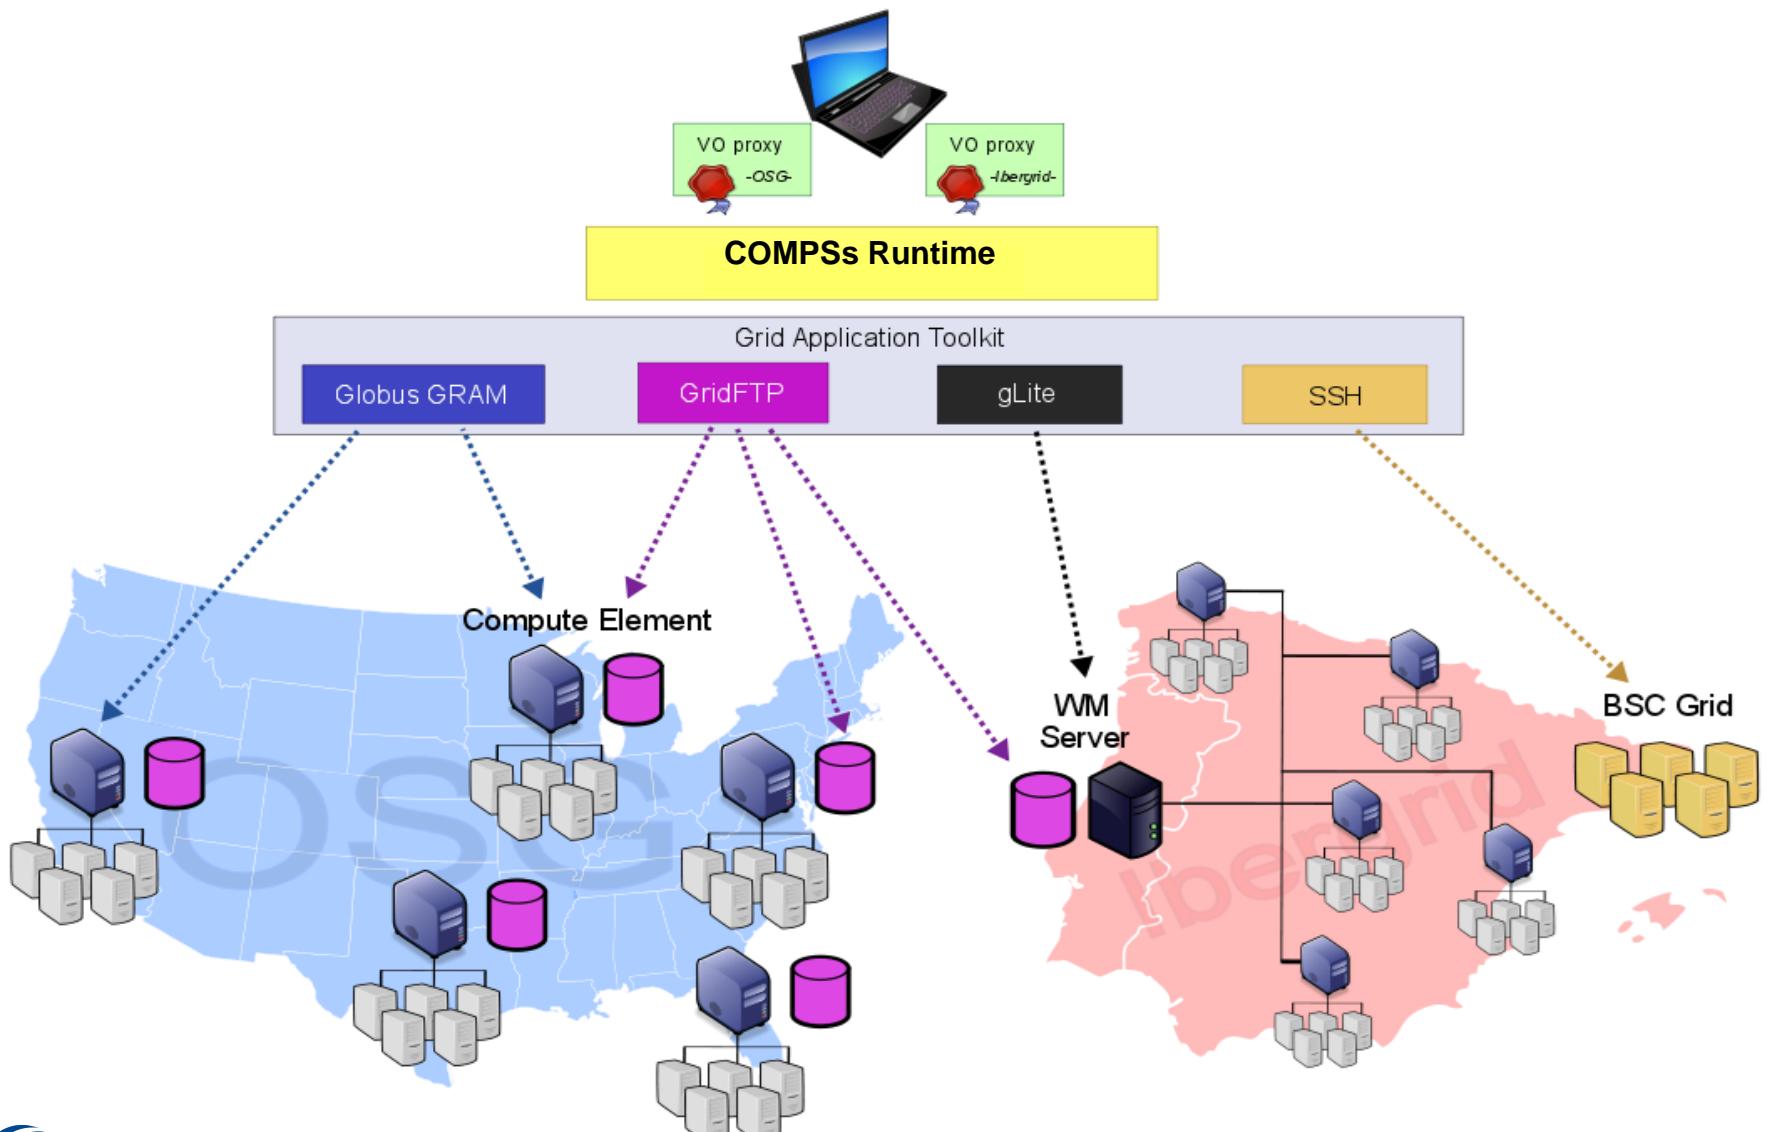


# Cloud Bursting

- Increase/decrease number of VMs depending on task load
- Bursting to Amazon EC2 to face peak load

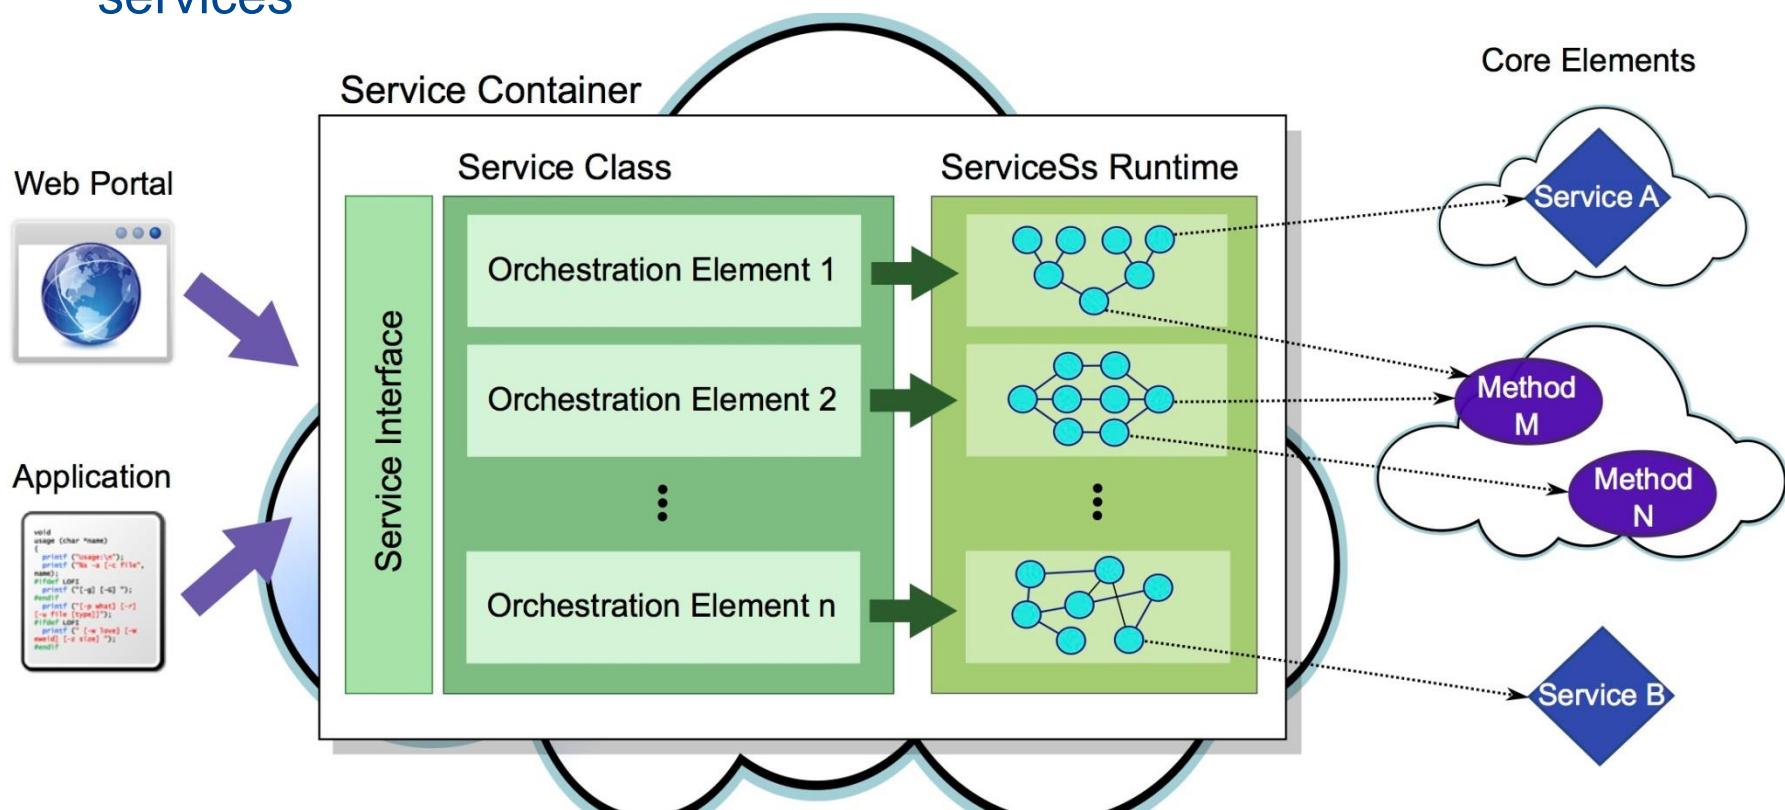


# COMPSs in a Grid



# Web service implementations with COMPSs

- « A WS method implements a workflow of tasks
- « Different invocations generate different tasks
- « Runtime manages the execution of the different calls in the available services



# IDE for COMPSs applications

## « IDE for implementing and deploying applications

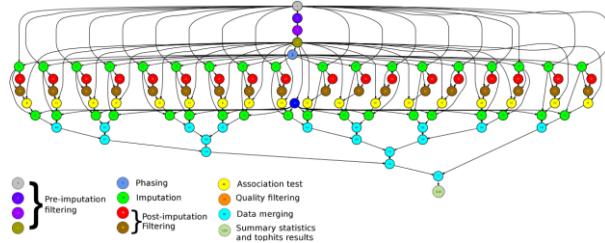
**Tasks Definition:**

- Service Operations (Orchestration)
- Tasks (Core Element)

**Building & Deployment:**

- Generate Packages
- Define hosts & Deploy

# COMPSs Framework



## Application repository

AlyaADAN

Gipsy/LOF  
AR

BioVeL

transPLAN  
T



EGI AppDB

## Enactment Service (PMES)

Programmatic interface to execute COMPSs applications

- OGSA-BES SOAP API
- Support to OCCI and CDMI
- Multiple cloud providers

### Web Dashboard

- Management of applications, jobs, storage and users in the PMES
- Monitoring of execution

## Programming Model

- Sequential programming model
- General purpose programming languages + annotations/hints
- Exploitation of implicit parallelism



- Automatic on-the-fly creation of a task dependency graph
- COMPSs workflows portable to HPC and Cloud without change



## Runtime

- Parallelization of task execution
- Hybrid executions (Private+Public)
- Elastic management of resources
- Interoperability through standards
- Automatic selection of VM templates depending on the task.



# Applications using COMPSs

## « Personalized medicine

- eIMRT: planning radiotherapy treatments

## « Earth Science

- HRT: modeling global ocean-atmosphere circulation

## « 3D render

- LuxRender: renderize architectural designs

## « Civil Engineering

- EnergyPlus: modeling airflows in buildings
- Architrave: force-effects on buildings

## « BioInformatics

- Discrete: simulate molecular dynamics for proteins
- Blast: alignments of protein sequences
- Hmmer: alignment of protein sequences
- GeneDetect: genetics algorithm
- QSAR: drug design



**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

# HANDS-ON

# Hands-On: Overview

## « COMPSs Virtual Machine Set-up

## « Java Hands-on

- Compilation & Execution
- Configuration
- Monitoring, debugging, graph generation

## « Python Hands-on

- Annotate tasks
- Execution in VM
- Overview of tracing, trace analysis
- Code optimization

# COMPSS development VM Installation

## « COMPSS Virtual Appliance

- Available from website
  - <http://compss.bsc.es/releases/vms/COMPSS-1.3-VM.ova>

## « VirtualBox: Import Virtual Appliance...

- user: compss, password: compss2016





**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

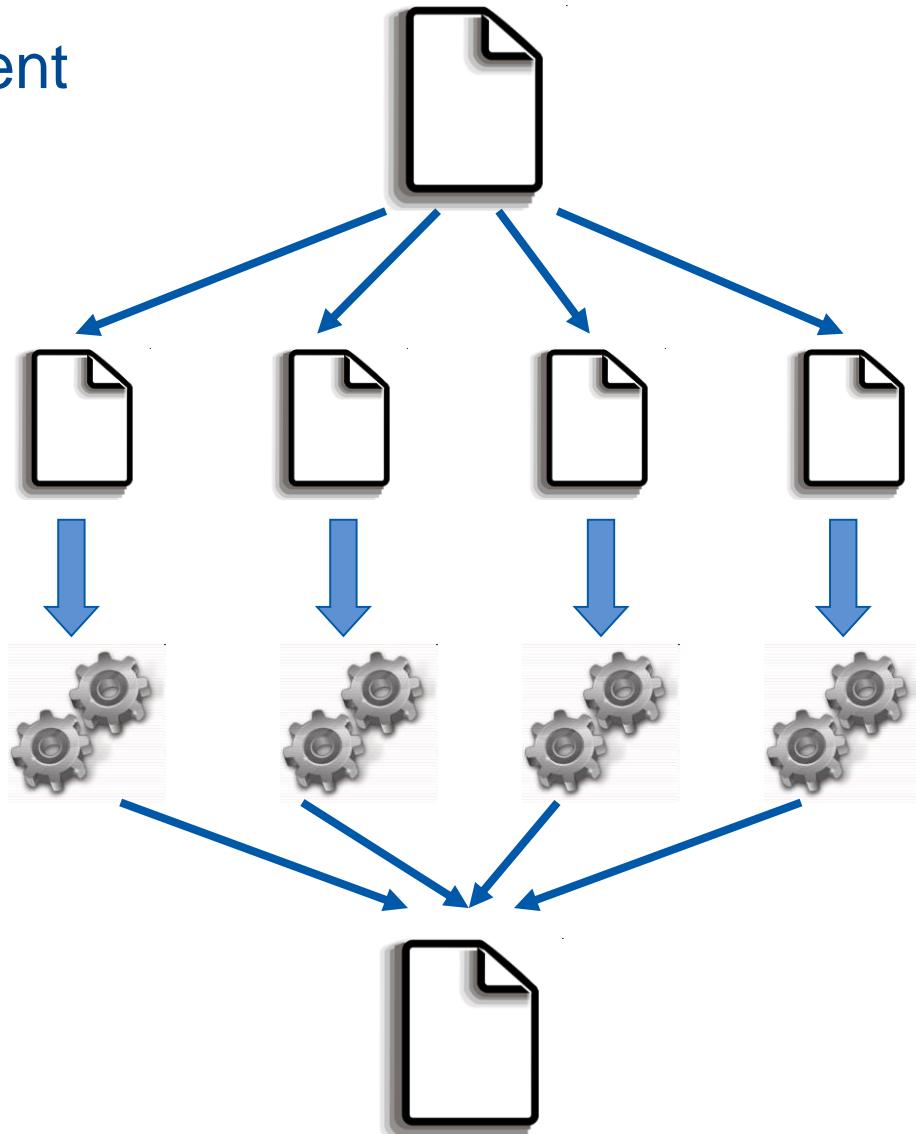
# Java Hands-on

# Word Count

« Counting words of a document

« Parallelization

- Split documents in blocks
- Count words of Blocks
- Merge results



# Java Hands On: Exercise

## « Complete the Word Count parallelization with COMPSS

- Level 0: No Java background
  - Look the implementation (wordcount project)
- Level 1: Basic Java background
  - Define methods in the interface (wordcount\_sequential)
- Level 2: Java background
  - Define methods in the interface and complete the part of the main code with helper methods (wordcount\_blanks)



# Java Hands On: Exercise Solution

## « Main Code

```
private static void computeWordCount() {
    HashMap<String, Integer> result = new HashMap<String, Integer>();
    int start = 0;
    for (int i = 0; i < NUM_BLOCKS; ++i) {
        HashMap<String, Integer> partialResult = wordCountBlock(DATA_FILE, start, BLOCK_SIZE);
        start = start + BLOCK_SIZE;
        result = mergeResults(result, partialResult);
    }
    System.out.println("[LOG] Counted Words is : " + result.keySet().size());
}
```

## « Interface

```
public interface WordcountItf {
    @Method(declaringClass = "wordcount.uniqueFile.Wordcount")
    public HashMap<String, Integer> mergeResults(
        @Parameter HashMap<String, Integer> m1,
        @Parameter HashMap<String, Integer> m2
    );

    @Method(declaringClass = "wordcount.uniqueFile.Wordcount")
    HashMap<String, Integer> wordCountBlock(
        @Parameter(type = Type.FILE, direction = Direction.IN) String filePath,
        @Parameter int start,
        @Parameter int bsize
    );
}
```

# Java Hands-on: Compilation and Simple Execution

## « Compilation (Eclipse IDE)

- Package Explorer -> Project (wordcount) -> Export... (Solution)

## « Use runcompss command to run the application

- runcompss [options] < FQDN app. classname> <application args>

## « **Exercise:** Simple wordcount execution

- Usage:

```
wordcount.uniqueFile.Wordcount <data_file> <block_size>
```



```
$compss@bsc:~/> cd ~/workspace_java/wordcount/jar  
$compss@bsc:~/workspace_java/wordcount/jar/> runcompss wordcount.uniqueFile.Wordcount  
/home/compss/workspace_java/wordcount/data/file_short.txt 650
```

# Java Hands-on: Result

```
$compss@bsc:~/workspace_java/wordcount/jar/> runcompss wordcount.uniqueFile.Wordcount  
/workspace_java/wordcount/data/file_short.txt 500
```

Using default location for project file:

`/opt/COMPSS/Runtime/scripts/user/../../configuration/xml/projects/project.xml`

Using default location for resources file:

`/opt/COMPSS/Runtime/scripts/user/../../configuration/xml/resources/resources.xml`

----- Executing `wordcount.uniqueFile.Wordcount` -----

WARNING: IT Properties file is null. Setting default values

[ API] - Deploying COMPSS Runtime v1.3 (build xxxx)

[ API] - Starting COMPSS Runtime v1.3 (build xxxx)

DATA\_FILE parameter value = /home/compss/workspace\_java/wordcount/data/file\_short.txt

BLOCK\_SIZE parameter value = 650

[LOG] Computing word count result

[LOG] Counted Words is : 250

[ API] - No more tasks for app 1

[ API] - Getting Result Files 1

[ API] - Execution Finished

Application Logs

# Java Hands-on: Configuration

## Project.xml:

/opt/COMPSS/Runtime/configuration/xml/projects/project.xml

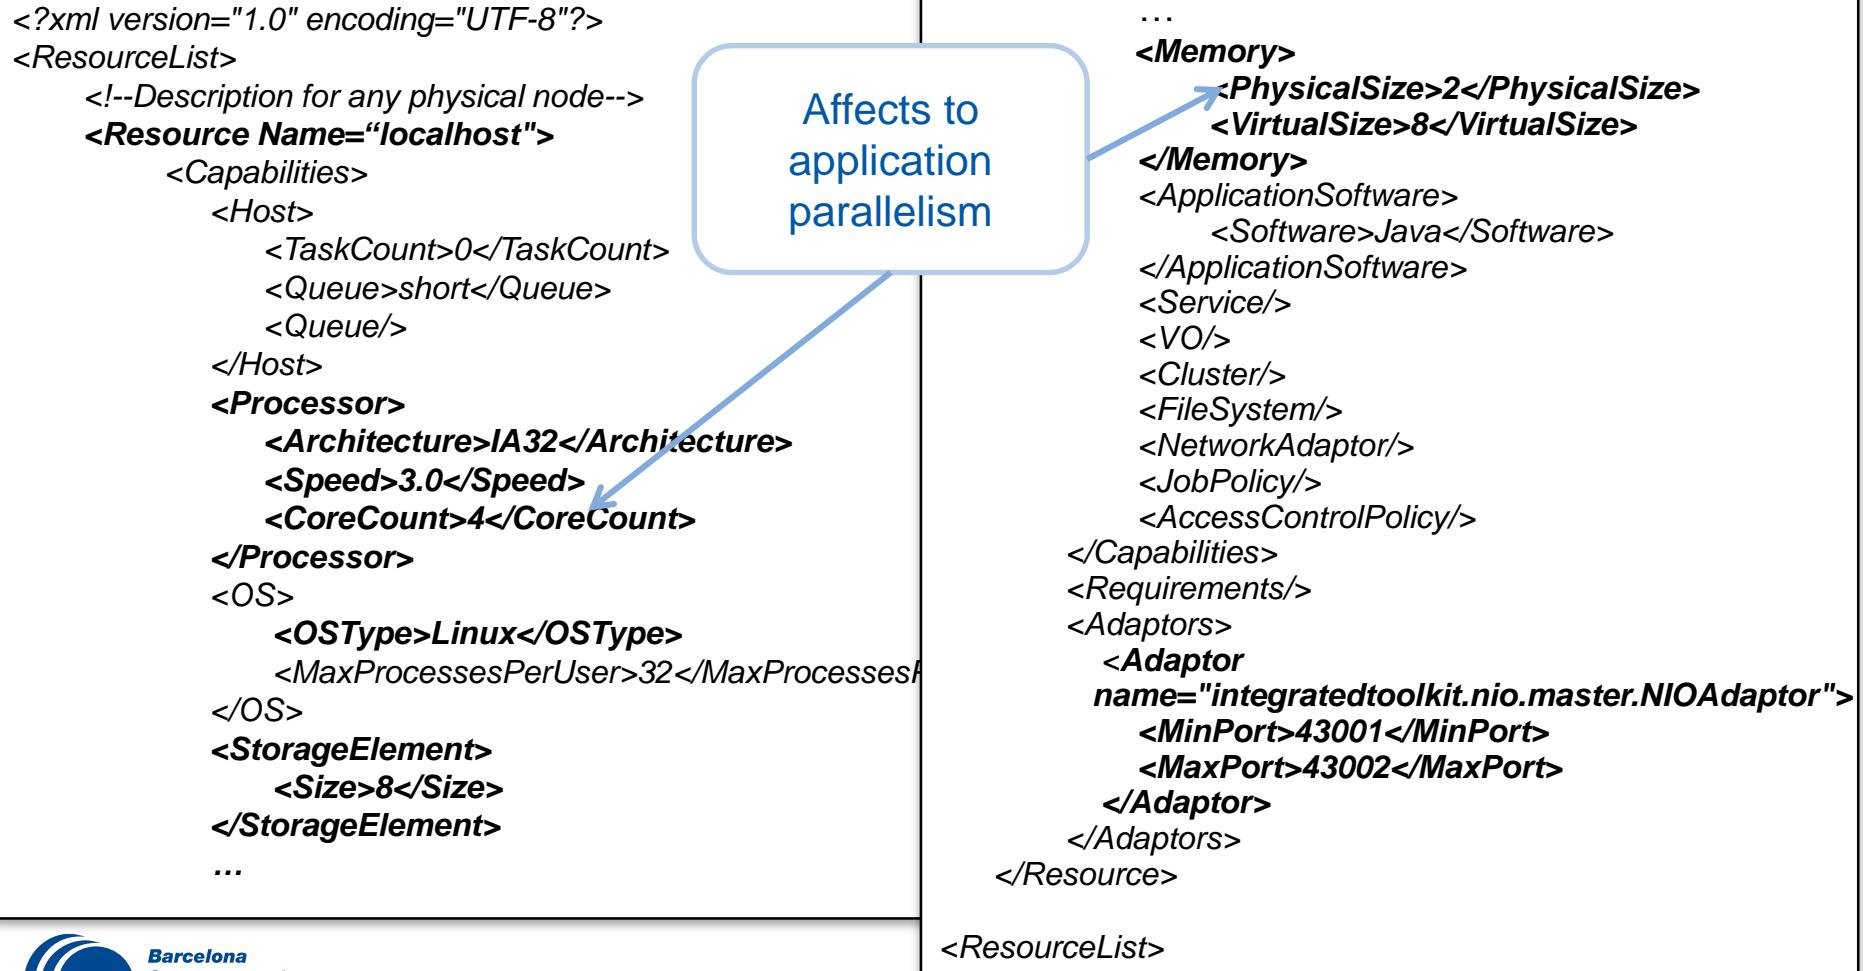
```
<?xml version="1.0" encoding="UTF-8"?>
<Project>
    <!--Description for any physical node-->
    <Worker Name="localhost">
        <InstallDir>/opt/COMPSS/Runtime/scripts/system</InstallDir>
        <WorkingDir>/tmp/localhost</WorkingDir>
    </Worker>
</Project>
```

- Other optional parameters
  - User, AppDir, LibraryPath

# Java Hands-On: Configuration

## Resources.xml:

/opt/COMPSSs/Runtime/configuration/xml/resources/resources.xml



# Java Hands-On: Monitoring

## « The runtime of COMPSS provides real-time monitoring

- <http://localhost:8080/compss-monitor>
- If not started run as root:
  - /etc/init.d/compss-monitor start

## « The user can log-in and follow the progress of the executions

- Running tasks, resources usage, execution time per task, real-time execution graph, etc.

## « Activate monitoring with a runcompss flag

- Setting a monitoring interval
  - runcompss **--monitoring=<int>**
- With a default monitoring interval
  - runcompss **-m** (or) runcompss **--monitoring**

## « **Exercise:** run wordcount enabling monitoring

```
$compss@bsc:~/> cd ~/workspace_java/wordcount/jar  
$compss@bsc:~/workspace_java/wordcount/jar/> runcompss -m wordcount.uniqueFile.Wordcount  
/home/compss/workspace_java/wordcount/data/file_long.txt 250000
```



# Java Hands-on: Debugging

- « Different log levels activated as runcompss options
  - runcompss **--log\_level=<level>**  
**(off**: for performance | **info**: basic logging | **debug**: detect errors)
  - runcompss **-debug** or runcompss **-d**
- « The output/errors of the main code of the application are shown in the console
- « Other logging files are stored in:
  - \$HOME/.COMPSs/<APP\_NAME>\_XX
- « Inside this folder, the user can check the following:
  - The output/error of a task # N : **/jobs/jobN.[out|err]**
  - Messages from the COMPSs : **runtime.log**
  - Task to resources allocation: **resources.log**
- « **Exercise:** run wordcount with debugging

```
$compss@bsc:~/> cd ~/workspace_java/wordcount/jar  
$compss@bsc:~/workspace_java/wordcount/jar/> runcompss -d wordcount.uniqueFile.Wordcount  
/home/compss/workspace_java/wordcount/data/file_short.txt 650
```



# Java Hands-on: Graph generation

- « To generate the graph of an application, it must be run with the monitor or graph flags activated
  - runcompss -m (or) runcompss –graph (or) runcompss -g
- « The graph will be stored in:
  - \$HOME/.COMPSs/<APP\_NAME>\_XX/monitor/complete\_graph.dot
- « To convert the graph to a PDF format use gengraph command
  - Usage: gengraph <dot\_file>
- « **Exercise:** generate the graph for the wordcount application



```
$compss@bsc:~/> cd ~/workspace_java/wordcount/jar  
$compss@bsc:~/workspace_java/wordcount/jar/> runcompss -g wordcount.uniqueFile.Wordcount  
/home/compss/workspace_java/wordcount/data/file_short.txt 650  
  
... application execution ...  
  
$compss@bsc:~/workspace_java/wordcount/jar/> cd ~/.COMPSs/wordcount.uniqueFile.Wordcount_04/monitor  
$~/.COMPSs/wordcount.uniqueFile.Wordcount_04/monitor> gengraph complete_graph.dot  
Output file: complete_graph.pdf  
$~/.COMPSs/wordcount.uniqueFile.Wordcount_04/monitor> evince complete_graph.pdf
```



**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

# Python Hands-on

# PyCOMPSs Hands On: Exercise

## Complete the WordCount parallelization with PyCOMPSs

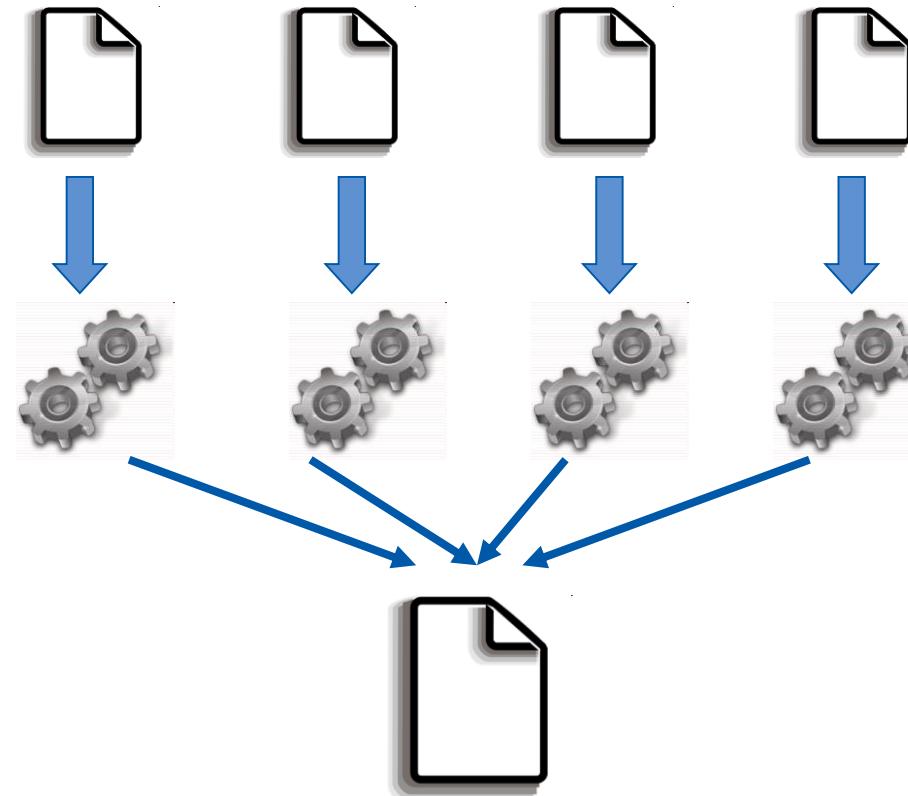
- Task definition with python decorators
- Execution in MareNostrum III
- Optimize the reduce method



# Word Count

- « Counting words of a set of documents
- « Parallelization

- Phase 1: Count words of a set of document
- Phase 2 : Merge results



# First step - Connecting to MareNostrum III

## How to connect to MN3?

- > ssh -X nct01XXX@mn3.bsc.es
- Password: COMPSS-0216.0XX



## Update .bashrc

- Edit: **.bashrc**
- Add: “**module load COMPSS/1.3**” at the end
- Execute: **source .bashrc**

## Where is the source code?

- cd**
- cp /gpfs/projects/nct01/nct01001/source/\* .**

launch\_pycompss.sh  
wordcount.py

## Where is the dataset?

- cp -r /gpfs/projects/nct01/nct01001/dataset/ .**

## Available editors

- vi
- emacs

# WordCount @ Sequential (wordcount.py)

```
def read_file(file_path):
    data = []
    with open(file_path, 'r') as f:
        for line in f:
            data += line.split()
    return data
```

```
def wordCount(data):
    partialResult = {}
    for entry in data:
        if entry in partialResult:
            partialResult[entry] += 1
        else:
            partialResult[entry] = 1
    return partialResult
```

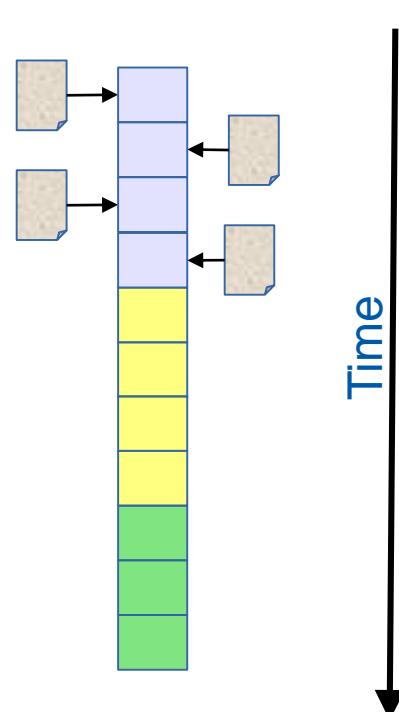
```
def merge_two_dicts(dic1, dic2):
    for k in dic2:
        if k in dic1:
            dic1[k] += dic2[k]
        else:
            dic1[k] = dic2[k]
    return dic1
```

```
if __name__ == "__main__":
    pathDataset = sys.argv[1]
    # Construct a list with the file's paths from the dataset
    paths = []
    for fileName in os.listdir(pathDataset):
        paths.append(os.path.join(pathDataset, fileName))

    # Read file's content
    data = map(read_file, paths)

    # From all file's data execute a wordcount on it
    partialResult = map(wordCount, data)

    # Accumulate the partial results to get the final result.
    result = reduce(merge_two_dicts, partialResult)
```



# WordCount @ Sequential

Remember the dataset path

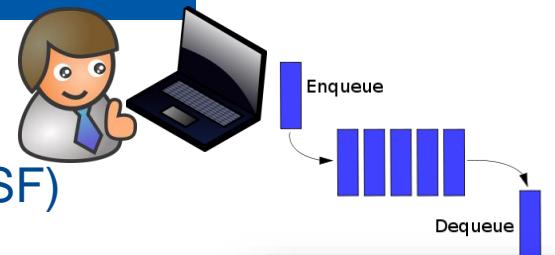
How to launch with python sequentially?

- > **python wordcount.py /home/nct01/nct01XXX/dataset/4files/**  
Results:

```
user@login3:~> python wordcount.py /path/to/dataset/
Elapsed Time (s)
0.959941864014
Words: 2545211
```

Submit jobs to MareNostrum III

- All jobs should be submitted to the queuing system (LSF)
- We will use a launcher script: **launch\_pycompss.sh**
- Useful commands:
  - bjobs – This command shows the status of the job.
  - bkill jobId – This command kills a job with id ‘jobId’.



# PyCOMPSs Hands On: Exercise

» Complete the WordCount parallelization with PyCOMPSs

- **Task definition with python decorators**
  - Modify the wordcount.py
- **Execution in MareNostrum III**
  - Use launch\_pycompss.sh
- Optimize the reduce method



# Execution in MareNostrum III - HandsOn

## «launch\_pycompss.sh

```
#!/bin/bash

enqueue_compss \
--exec_time=10 \
--num_nodes=2 \
--tasks_per_node=8 \
--lang=python \
--classpath=/gpfs/home/nct01/nct01XXX/ \
--tracing=true \
--graph=true \
/home/nct01/nct01XXX/wordcount.py /gpfs/home/nct01/nct01XXX/dataset/64files
```

## «Parameters:

- num\_nodes: amount of nodes where to execute (1 master + 1 worker).
- tasks\_per\_node: amount of tasks that can be processed in parallel (1-16).
- Dataset path: **/gpfs/home/nct01/nct01XXX/dataset/64files**

## «How to execute with PyCOMPSs?

- **./launch\_pycompss.sh**

# WordCount @ PyCOMPSSs (option1)

```
def read_file(file_path):
    data = []
    with open(file_path, 'r') as f:
        for line in f:
            data += line.split()
    return data
```

```
@task(returns=dict)
def wordCount(data):
    partialResult = {}
    for entry in data:
        if entry in partialResult:
            partialResult[entry] += 1
        else:
            partialResult[entry] = 1
    return partialResult
```

```
@task(returns=dict, priority=True)
def merge_two_dicts(dic1, dic2):
    for k in dic2:
        if k in dic1:
            dic1[k] += dic2[k]
        else:
            dic1[k] = dic2[k]
    return dic1
```

```
from pycompss.api.task import task
from pycompss.api.parameter import *

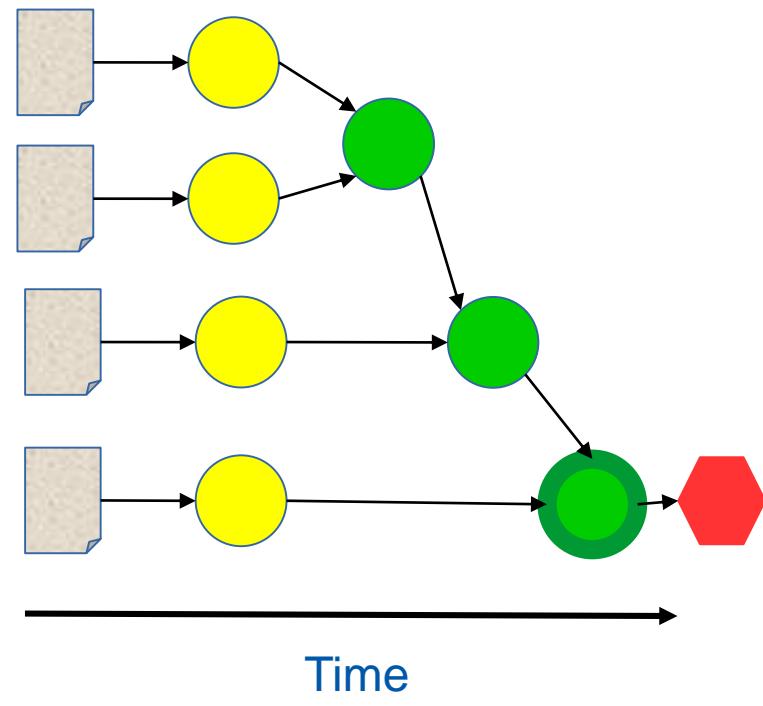
if __name__ == "__main__":
    from pycompss.api.api import compss_wait_on
    pathDataset = sys.argv[1]
    # Construct a list with the file's paths from the dataset
    paths = []
    for fileName in os.listdir(pathDataset):
        paths.append(os.path.join(pathDataset, fileName))

    # Read file's content
    data = map(read_file, paths)

    # From all file's data execute a wordcount on it
    partialResult = map(wordCount, data)

    # Accumulate the partial results to get the final result.
    result = reduce(merge_two_dicts, partialResult)

    # Wait for result
    result = compss_wait_on(result)
```



# WordCount @ PyCOMPSs (option 2)

```
@task(returns=list,
      file_path=FILE_in)
def read_file(file_path):
    data = []
    with open(file_path, 'r') as f:
        for line in f:
            data += line.split()
    return data
```

```
@task(returns=dict)
def wordCount(data):
    partialResult = {}
    for entry in data:
        if entry in partialResult:
            partialResult[entry] += 1
        else:
            partialResult[entry] = 1
    return partialResult
```

```
@task(returns=dict, priority=True)
def merge_two_dicts(dic1, dic2):
    for k in dic2:
        if k in dic1:
            dic1[k] += dic2[k]
        else:
            dic1[k] = dic2[k]
    return dic1
```

```
from pycompss.api.task import task
from pycompss.api.parameter import *

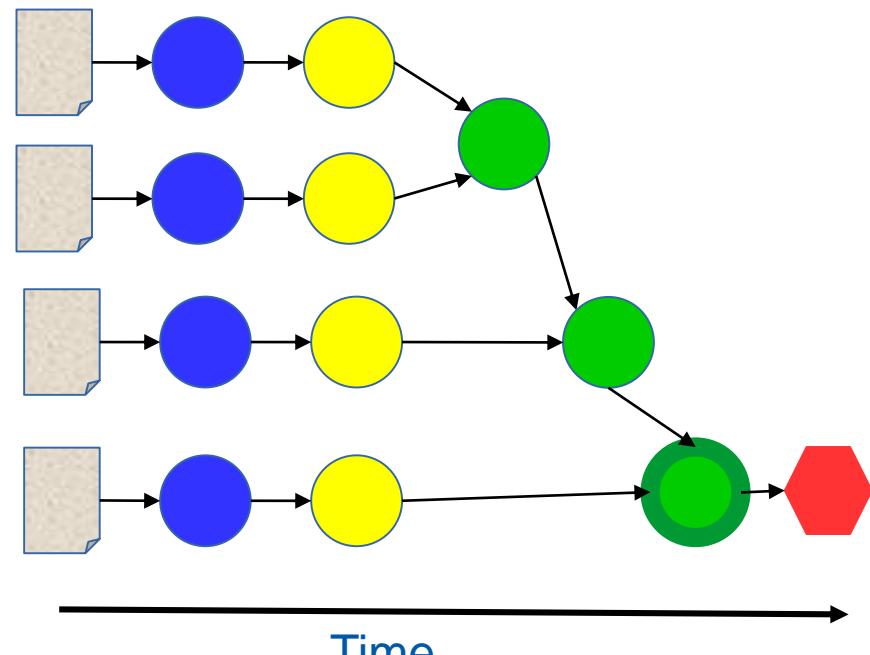
if __name__ == "__main__":
    from pycompss.api.api import compss_wait_on
    pathDataset = sys.argv[1]
    # Construct a list with the file's paths from the dataset
    paths = []
    for fileName in os.listdir(pathDataset):
        paths.append(os.path.join(pathDataset, fileName))

    # Read file's content
    data = map(read_file, paths)

    # From all file's data execute a wordcount on it
    partialResult = map(wordCount, data)

    # Accumulate the partial results to get the final result.
    result = reduce(merge_two_dicts, partialResult)

    # Wait for result
    result = compss_wait_on(result)
```



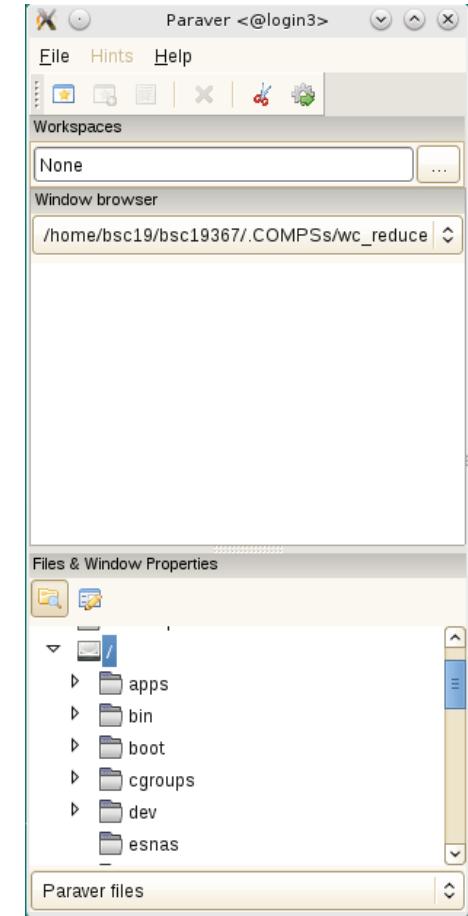
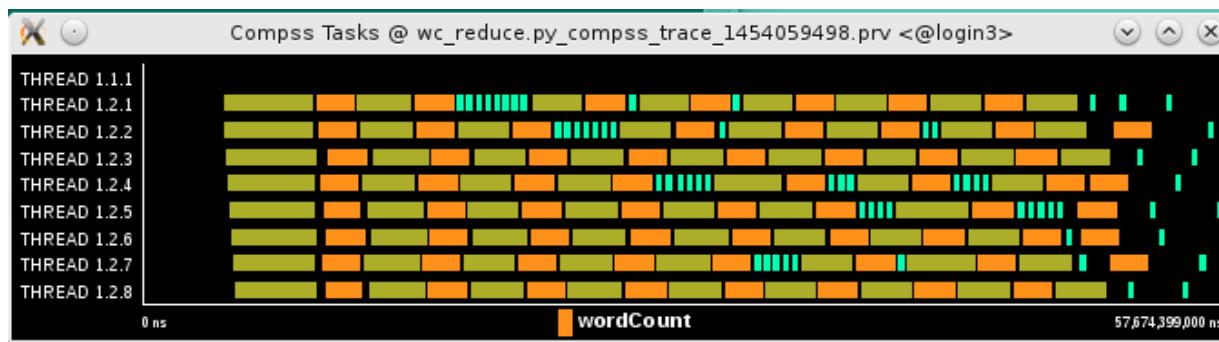
# Wordcount @ Performance Analysis

¶ Paraver is the BSC tool for trace visualization

- Trace events are encoded in Paraver (.prv) format by Exrae
- Paraver is a powerful tool for trace visualization.
- An experienced user could obtain many different views of the trace events.

¶ For more information about Paraver visit:

- <http://www.bsc.es/computer-sciences/performance-tools/paraver>



# Wordcount @ Performance Analysis

COMPSs can generate post-execution traces of the distributed execution of the application

- Useful for performance analysis and diagnosis

How it works?

- Task execution and file transfers are application events
- An XML file is created at workers to keep track of these events
- At the end of the execution all the XML files are merged to get the final trace file
- COMPSs uses Extrae tool to dynamically instrument the application
  - In a worker:
    - Extrae keeps track of the events in an intermediate file
  - In the master:
    - Extrae merges the intermediate files to get the final trace file

# Wordcount @ Performance Analysis

----- Executing wc\_reduce.py -----

Welcome to Extrae 3.1.1rc (revision 3360 based on extrae/trunk)

Extrae: Generating intermediate files for Paraver traces.

Extrae: Intermediate files will be stored in ./statelite/tmpfs/gpfs/home/bsc19/bsc19000/Apps/WC/src/tutorial

Extrae: Tracing buffer can hold 500000 events

Extrae: Tracing mode is set to: Detail.

Extrae: Successfully initiated with 1 tasks

[ API] - Deploying COMPSs Runtime v1.3 (build 20151211-1532)

[ API] - Starting COMPSs Runtime v1.3 (build 20151211-1532)

...

[ API] - No more tasks for app 0

[ API] - Getting Result Files 0

[ API] - Execution Finished

...

Extrae starts before the user application execution

COMPSs runtime starts

The application finishes and the tracing process ends

COMPSs runtime ends

Extrae: Application has ended. Tracing has been terminated.

The merge process starts

merger: Output trace format is: Paraver

merger: Extrae 3.1.1rc (revision 3360 based on extrae/trunk)

Intermediate trace files are processed

mpi2prv: Selected output trace format is Paraver

mpi2prv: Parsing intermediate files

mpi2prv: Generating tracefile (intermediate buffers of 745642 events)

The final trace file is generated

# WordCount @ Performance Analysis

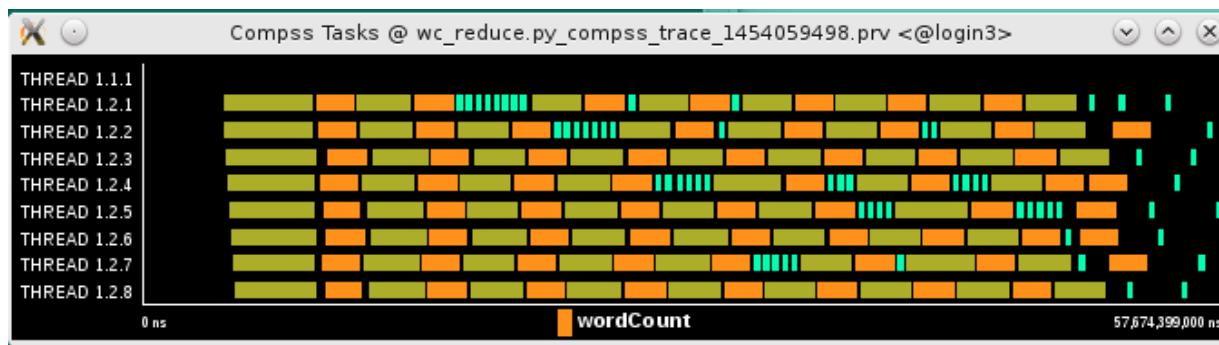
## « Open Paraver

- > module load paraver
- > cd \$HOME/.COMPSSs/wordcount.py\_01
- > wxparaver trace/\*.prv

## « COMPSS provides some configuration files to automatically obtain the view of the trace

### □ File/Load Configuration:

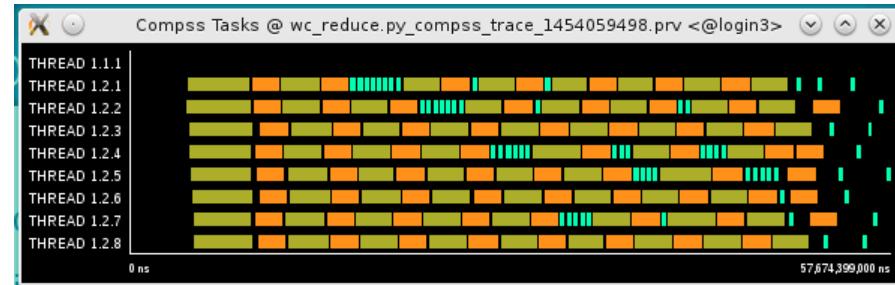
(/gpfs/apps/MN3/COMPSSs/1.3/Dependencies/paraver/cfgs/compss\_tasks.cfg)



# Wordcount @ Performance Analysis

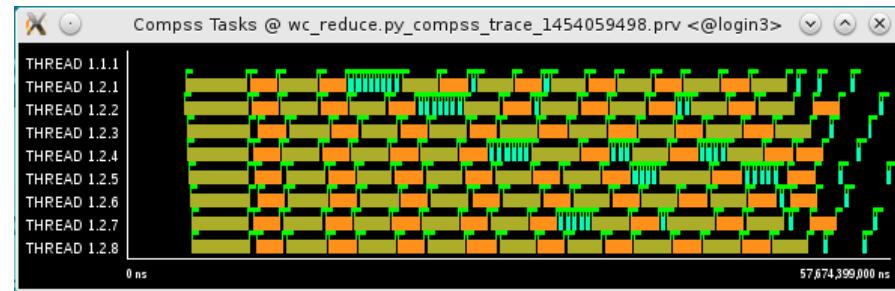
## Fit window

- Right click on the trace window
- Fit Semantic Scale/ Fit Both



## View Event flags

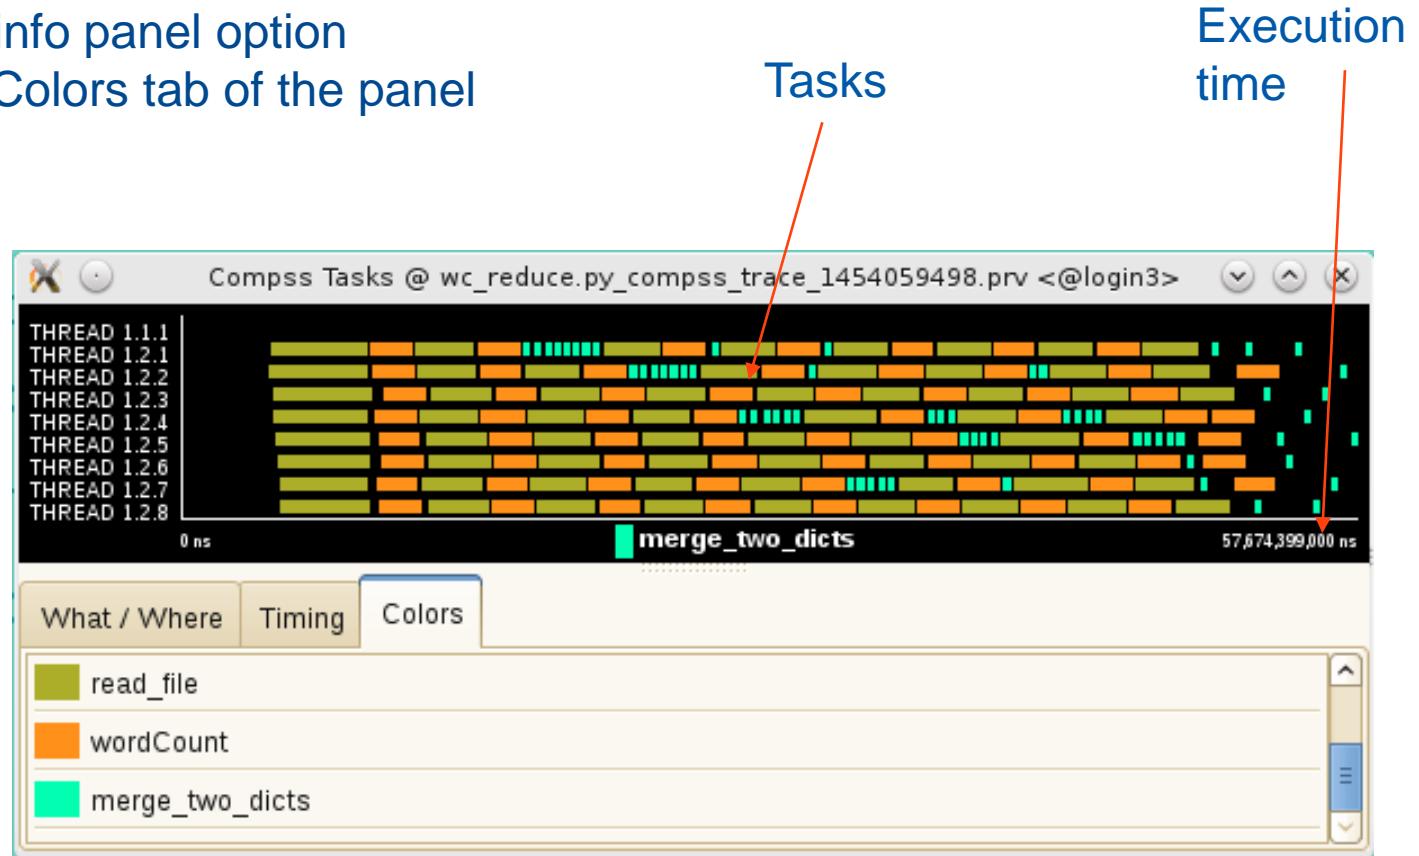
- Right click on the trace window
- View / Event Flags



# Wordcount @ Performance Analysis

## >Show info Panel

- Right click on the trace window
- Check info panel option
- Select Colors tab of the panel



# Wordcount @ Performance Analysis

## Zoom to see details

- Select a region in the trace window to see in detail
- And repeat the process until the needed zoom level
- The undo zoom option is in the right click panel



Previous task  
ends

Processor is  
idle

New task starts

# Wordcount @ Performance Analysis

## Summarizing:

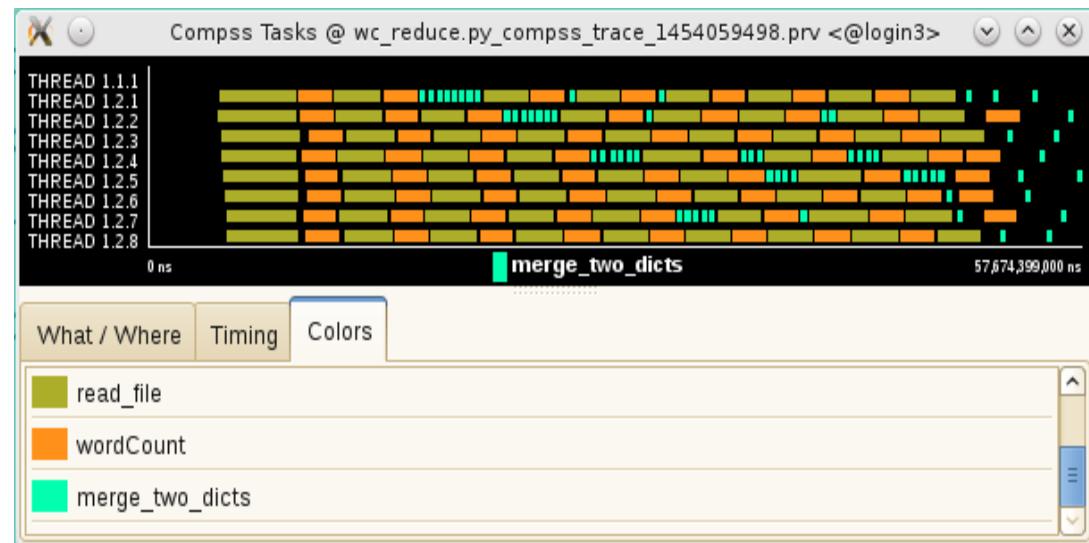
- Lines in the trace:
  - One line for the master
  - N lines for the workers

## Meaning of the colours:

- Black: idle
- Other colors: task running
  - see the color legend

## Flags (events):

- Start / end of task



# PyCOMPSs Hands On: Exercise

## Complete the WordCount parallelization with PyCOMPSs

- Task definition with python decorators
- Execution in MareNostrum III
- Optimize the reduce method



# WordCount @ PyCOMPSs (option 2)

```
@task(returns=list,
      file_path=FILE_in)
def read_file(file_path):
    data = []
    with open(file_path, 'r') as f:
        for line in f:
            data += line.split()
    return data
```

```
@task(returns=dict)
def wordCount(data):
    partialResult = {}
    for entry in data:
        if entry in partialResult:
            partialResult[entry] += 1
        else:
            partialResult[entry] = 1
    return partialResult
```

```
@task(returns=dict, priority=True)
def merge_two_dicts(dic1, dic2):
    for k in dic2:
        if k in dic1:
            dic1[k] += dic2[k]
        else:
            dic1[k] = dic2[k]
    return dic1
```

```
from pycompss.api.task import task
from pycompss.api.parameter import *

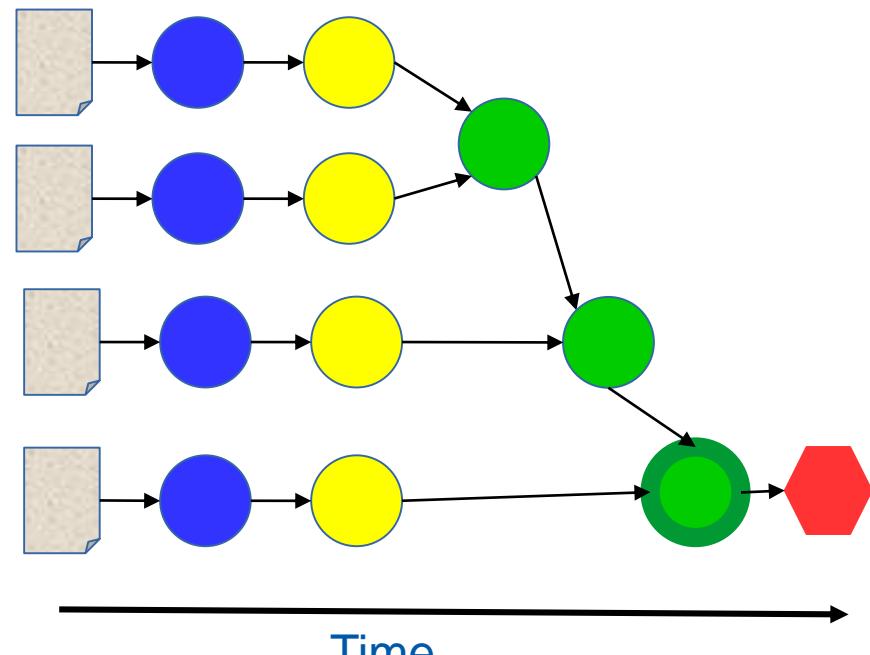
if __name__ == "__main__":
    from pycompss.api.api import compss_wait_on
    pathDataset = sys.argv[1]
    # Construct a list with the file's paths from the dataset
    paths = []
    for fileName in os.listdir(pathDataset):
        paths.append(os.path.join(pathDataset, fileName))

    # Read file's content
    data = map(read_file, paths)

    # From all file's data execute a wordcount on it
    partialResult = map(wordCount, data)

    # Accumulate the partial results to get the final result.
    result = reduce(merge_two_dicts, partialResult)

    # Wait for result
    result = compss_wait_on(result)
```



# WordCount @ PyCOMPSs (Optimization)

```
@task(returns=list,
      file_path=FILE_IN)
def read_file(file_path):
    data = []
    with open(file_path, 'r') as f:
        for line in f:
            data += line.split()
    return data
```

```
from pycompss.api.task import task
from pycompss.api.parameter import *

if __name__ == "__main__":
    from pycompss.api.api import compss_wait_on
    pathDataset = sys.argv[1]
    # Construct a list with the file's paths from the dataset
    paths = []
    for fileName in os.listdir(pathDataset):
        paths.append(os.path.join(pathDataset, fileName))

    # Read file's content
    data = map(read_file, paths)

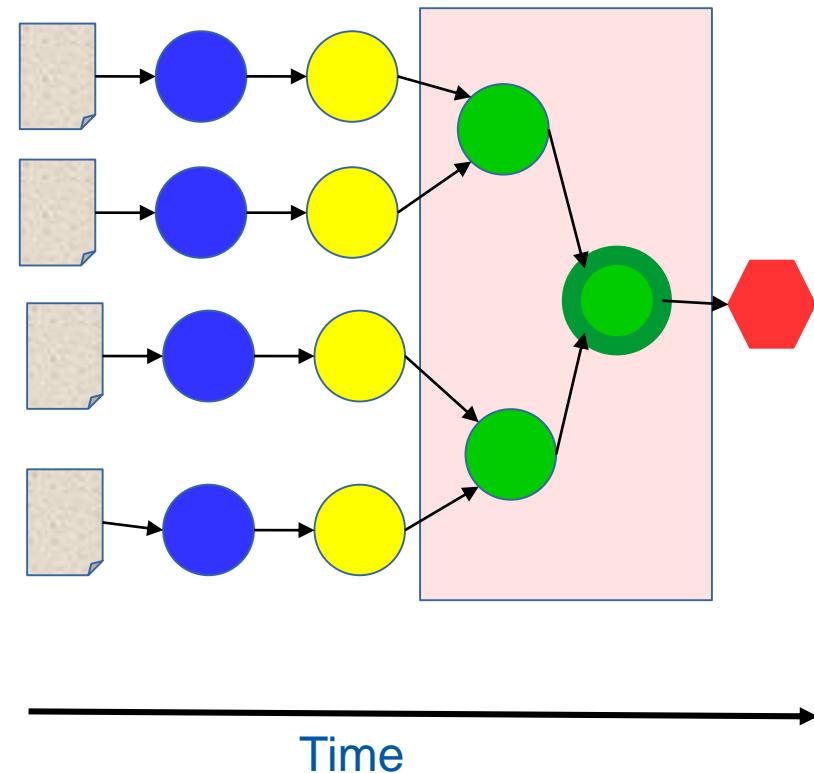
    # From all file's data execute a wordcount on it
    partialResult = map(wordCount, data)

    # Accumulate the partial results to get the final result.
    result = mergeReduce(merge_two_dicts, partialResult)

    # Wait for result
    result = compss_wait_on(result)
```

```
@task(returns=dict)
def wordCount(data):
    partialResult = {}
    for entry in data:
        if entry in partialResult:
            partialResult[entry] += 1
        else:
            partialResult[entry] = 1
    return partialResult
```

```
@task(returns=dict, priority=True)
def merge_two_dicts(dic1, dic2):
    for k in dic2:
        if k in dic1:
            dic1[k] += dic2[k]
        else:
            dic1[k] = dic2[k]
    return dic1
```



# Execution in MareNostrum III - HandsOn

## «launch\_pycompss.sh

```
#!/bin/bash

enqueue_compss \
--exec_time=10 \
--num_nodes=2 \
--tasks_per_node=8 \
--lang=python \
--classpath=/gpfs/home/nct01/nct01XXX/ \
--tracing=true \
--graph=true \
/home/nct01/nct01XXX/wordcount_optimized.py /gpfs/home/nct01/nct01XXX/dataset/64files
```

## «Parameters:

- num\_nodes: amount of nodes where to execute (1 master + 1 worker).
- tasks\_per\_node: amount of tasks that can be processed in parallel (1-16).
- Dataset path: **/gpfs/home/nct01/nct01XXX/dataset/64files**

## «How to execute with PyCOMPSs?

- **./launch\_pycompss.sh**

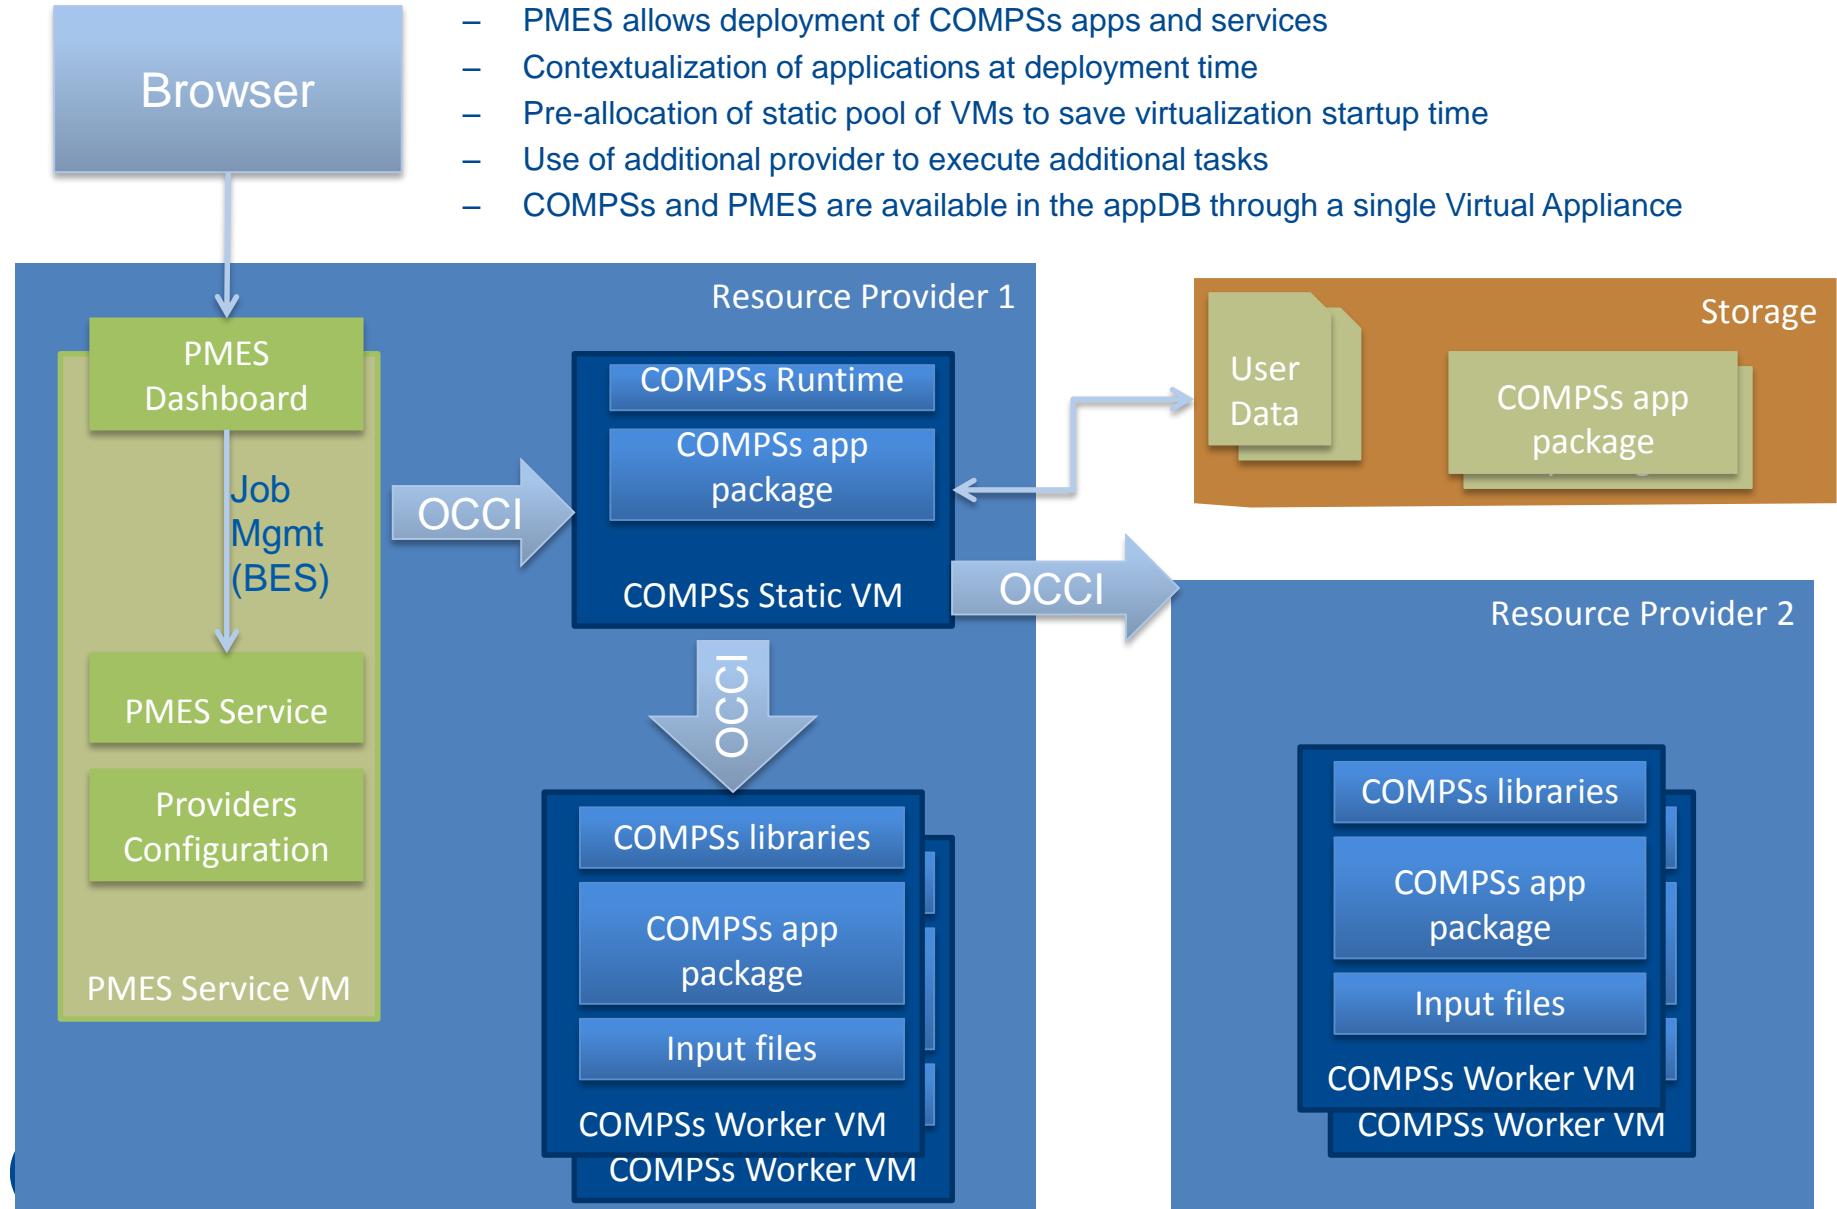


**Barcelona  
Supercomputing  
Center**

Centro Nacional de Supercomputación

# COMPSs and EGI HANDS-ON

# Deployment of COMPSSs and PMES on EGI



# Configuration

## « Choose the fedcloud.egi.eu VO

- UPV Provider: <https://fc-one.i3m.upv.es:11443>
- Pre Instantiated VMs

## « EGI AppDB Image

- `uuid_image_for_compss_pmes_ubuntu1504kv_im153_128`
- User: compss Password: compss2015

## « COMPSs configuration

- Configuration files for local and cloud execution available at:
  - <http://bscgrid05.bsc.es/~leaggi/local.tgz>
  - <http://bscgrid05.bsc.es/~leaggi/cloud.tgz>

## « Prepare the COMPSs package

- Blast application in Java workspace

## « Monitoring

- <http://158.42.105.X:8080/compss-monitor>

# Final Notes

- « Sequential programming approach
- « Parallelization at task level
- « Transparent data management and remote execution
- « Can operate on different infrastructures:
  - Cluster
  - Grid
  - Cloud (Public/Private)
    - PaaS
    - IaaS
  - Web services

# Final Notes

## « Project page:

- <http://www.bsc.es/compss>

## « Direct downloads page:

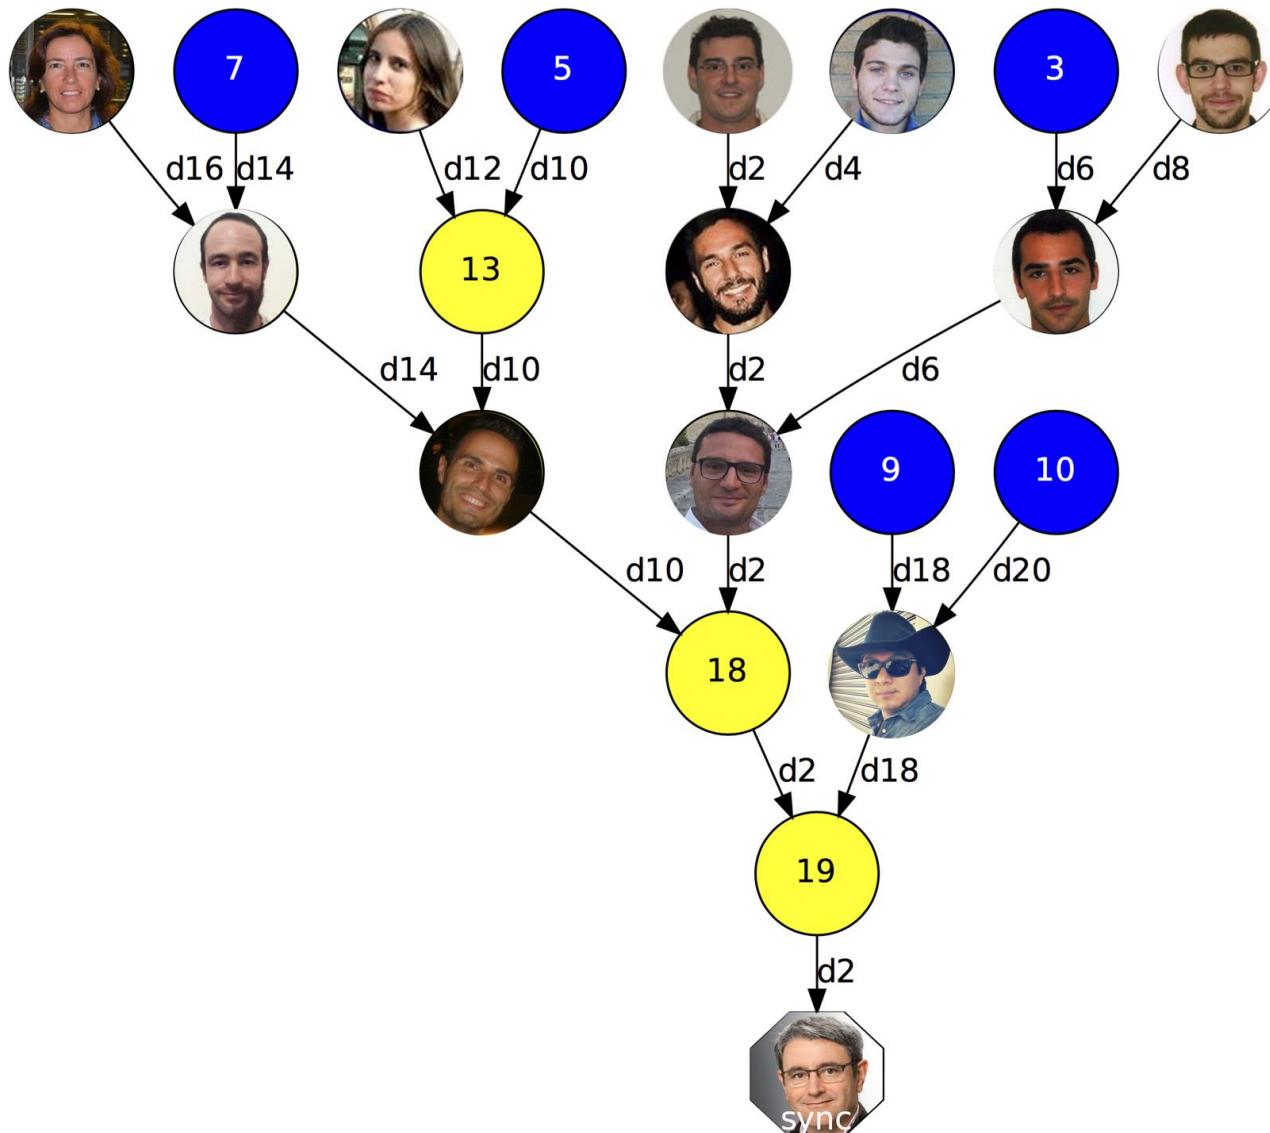
- <http://www.bsc.es/computer-sciences/grid-computing/comp-superscalar/download>

- Virtual Appliance for testing & sample applications
- Tutorials
- Red-Hat & Debian based installation packages
- Source Code

## « Application Repository

- <http://compss.bsc.es/projects/bar/wiki/Applications>
- Several examples of applications developed with COMPSs

# COMPSS Group





***Barcelona  
Supercomputing  
Center***  
*Centro Nacional de Supercomputación*

**Thank you!**

For further information please contact

[support-compss@bsc.es](mailto:support-compss@bsc.es)